

# Calibration Versus Scene Switching of Detector Operating Points for Aerial Multi-Object Tracking

Ahmed Gouda Ismail\* and Mohamed S. Mohamed<sup>†</sup>  
*Military Technical College, Cairo, Egypt*

Tarek Ahmed Mahmoud<sup>‡</sup>  
*Egypt University of Informatics, Cairo, Egypt*

Multi-object tracking from unmanned aerial vehicles must cope with scenes that change in density, object scale, illumination, and sharpness within a fixed compute budget. This paper asks how much of the benefit of adapting the detector operating point (confidence threshold, suppression threshold, and input resolution) comes from switching it with the scene and how much from calibrating it. The instrument is a causal, training-free controller driven by a scene complexity index of five cues ranked against a label-free calibration set; its eleven parameters were selected on labelled validation data and frozen before a locked held-out test. On seventeen held-out aerial sequences, the controller raised higher order tracking accuracy from 28.43 to 33.84 percent over a static default at 38.98 frames per second on one Tesla T4 graphics processor, excluding image decoding, and it transferred to a second aerial benchmark without retuning (24.08 to 28.39 percent). A matched static configuration reached 33.93 percent, so the gain comes from the validation-calibrated operating point and tracker profile, not from frame-to-frame switching. A separately recalibrated controller on a second pipeline showed no measurable held-out change against that tracker’s published static setting; a third pipeline gained slightly in-sample, but not over its default static setting.

## Nomenclature

|             |   |                                                                                 |
|-------------|---|---------------------------------------------------------------------------------|
| $c_t$       | = | detector confidence threshold applied at frame $t$                              |
| $c_e, c_h$  | = | confidence threshold at $S = 0$ (easy end) and at $S = 1$ (hard end)            |
| $\hat{F}_k$ | = | empirical distribution function of cue $k$ on the calibration set               |
| $K$         | = | analysis stride, frames between two cue analyses                                |
| $N_t$       | = | number of detections kept by the controller-set confidence filter at frame $t$  |
| $q_k$       | = | normalized value of scene cue $k \in \{\text{crowd, tiny, edge, light, blur}\}$ |

---

\*Computer Engineering and Artificial Intelligence Department; corresponding author, a7medgouda1@gmail.com.

<sup>†</sup>Computer Engineering and Artificial Intelligence Department, mohamedms@mtc.edu.eg.

<sup>‡</sup>Computer Engineering Department, tarek.mohamed@eui.edu.eg.

|                  |   |                                                                                     |
|------------------|---|-------------------------------------------------------------------------------------|
| $r_t$            | = | detector input resolution at frame $t$ , pixels                                     |
| $r_0, r_1, r_2$  | = | low, intermediate, and high resolution levels, pixels                               |
| $S_t$            | = | smoothed scene complexity index at frame $t$                                        |
| $s_j$            | = | raw scene complexity index of analysis $j$                                          |
| $W$              | = | smoothing window, number of analyses averaged                                       |
| $w_k$            | = | normalized weight of cue $k$                                                        |
| $\beta$          | = | mean gray level of the analysis image                                               |
| $\eta_t$         | = | overlap (intersection over union) threshold of non-maximum suppression at frame $t$ |
| $\eta_e, \eta_h$ | = | suppression threshold at $S = 0$ and at $S = 1$                                     |
| $\theta_t$       | = | detector operating point ( $c_t, \eta_t, r_t$ )                                     |
| $\tau_m, \tau_h$ | = | complexity thresholds that switch resolution to $r_1$ and to $r_2$                  |
| $\nu$            | = | variance of the Laplacian of the analysis image                                     |
| $\gamma$         | = | mean Sobel gradient magnitude of the analysis image                                 |
| $\Delta$         | = | paired difference between two systems, percentage points                            |

## I. Introduction

CAMERAS carried by small unmanned aerial vehicles are now routine sensors for traffic monitoring, inspection, search, and surveillance, and most of these missions need more than detection: they need persistent identities for the objects in view. Tracking by detection is the dominant way to obtain those identities in real time: a detector proposes boxes in each frame and an online tracker associates them over time [1–3]. Aerial video is a difficult case for this pipeline. Objects are small and numerous, the platform moves, and the scene changes within a single flight from sparse rural roads to dense intersections, from daylight to dusk, and from sharp to motion-blurred frames [4–6]. The same airborne perception problem appears throughout the aerospace information systems literature, from ground-target tracking by multirotor platforms [7] and compact descriptors for moving-object tracking [8] to deep detectors deployed on board [9, 10] and simulated imagery for vision-based refueling [11].

A detector deployed in such a pipeline is configured by a small number of operating-point parameters: the confidence threshold that decides which boxes are passed to the tracker, the overlap threshold of non-maximum suppression (NMS), and the input resolution. They are usually set once, often to library defaults, and left unchanged. That choice has two costs. First, the defaults were chosen for generic ground-level benchmarks, not for aerial footage with many objects of a few pixels. Second, a fixed setting cannot follow the scene: a threshold that suppresses clutter in a crowded intersection may discard the few visible vehicles on an empty road, and the resolution needed to see small objects is wasted on frames that contain none. Resolution also dominates the compute budget, which matters on a platform with a fixed

processor and a hard frame rate.

This paper studies whether a lightweight, causal controller that reads the scene and adjusts the detector operating point on the fly improves aerial multi-object tracking under a processing-throughput constraint of 25 frames per second, and, above all, whether any improvement comes from the adaptation itself or from the operating point that the adaptation settles on. Content-adaptive control is usually evaluated against a default configuration; this paper evaluates it against the matched static operating point as well. The controller combines five inexpensive cues, namely detection density, the fraction of tiny objects, image edge strength, low light, and blur, into a scalar scene complexity index (SCI) and maps that index to the three operating-point parameters. It learns nothing at run time; its cue calibration uses no labels, its parameters are selected with the labels of a small validation set, and its lightweight scene analysis is included in the measured processing throughput.

The contribution of the paper is a controlled study that separates the value of calibrating the detector operating point from the value of switching it with the scene, using a scene-adaptive controller as the instrument. The specific contributions are as follows:

- 1) An attribution analysis showing that the validation-optimized scene-adaptive controller performs like its own matched static operating point on held-out data, so that its gain over a static default and over a hand-designed controller is a calibration gain, together with ablation evidence on when scene-driven switching does and does not add accuracy (Secs. IX–XII).
- 2) A locked held-out comparison on the VisDrone2019 multi-object tracking test-dev split and a zero-tuning cross-dataset test on UAVDT, with paired sequence-level bootstrap intervals (Sec. IX), complemented by a cross-pipeline check on two further detector–tracker pipelines (Sec. XIII).
- 3) The instrument of the study: a training-free scene-adaptive controller whose cues are ranked against a label-free calibration set derived from fixed detector outputs, so that the cue scale does not depend on ground truth or on the test data (Sec. IV), and a validation-only selection chain of operating-point sweeps, a temporal ablation, and single- and multi-objective searches with predeclared selection rules (Secs. V–VIII).

In short, the optimized controller improves substantially over the static default and the hand-designed controller and generalizes to a second aerial dataset without retuning, but a matched static configuration, evaluated after the freeze for attribution, reaches the same held-out accuracy: the gain comes from the validation-calibrated operating point and tracker profile, and switching contributes no measurable accuracy for this detector and operating range. The paper explains why the optimization converged to near-static behavior and uses the hand-designed controller’s ablation to show where scene-driven resolution switching does pay off.

## II. Related Work

### A. Aerial detection and tracking

The VisDrone benchmark [4] and the UAVDT benchmark [5] established large annotated drone video collections for detection and multi-object tracking, and they document the characteristic difficulties of the domain: small objects, high density, changing altitude and viewpoint, and illumination changes. Surveys of deep learning for drone-based detection and tracking [6] and of computer vision for unmanned aerial vehicles [12] summarize the resulting methods. Methods specialized for moving aerial platforms model camera motion explicitly [13], and the aerospace literature has studied airborne target tracking with explicit motion models [7], efficient feature descriptors for ground moving-object tracking from small aircraft [8], onboard deep detection for real-time urban object detection [10], and onboard vision-based deep learning for docking [9]. Small-object detection, the central difficulty of aerial footage, is surveyed in [14]; multi-scale feature pyramids [15] are the standard architectural response, and input resolution remains the simplest operational lever.

### B. Tracking by detection

Online trackers of the SORT family combine a constant-velocity Kalman filter [16] with Hungarian assignment [17] on box overlap [1]; DeepSORT adds appearance features [2]; the IOU tracker shows how far pure overlap association can go at high frame rates [18]. ByteTrack associates high-confidence detections first and then recovers occluded objects from low-confidence detections [3], and observation-centric and appearance-augmented variants refine the motion model and the association cost [19–22]. Joint detection and embedding [23] and transformer-based end-to-end trackers [24] remove the separate association step. Nearly all of these methods take the detector’s operating point as given and tune their own thresholds against it, which is why the detector operating point and the tracker profile are treated together in this paper.

### C. Detector operating points

Single-stage detectors [25] and two-stage detectors [26] both expose a confidence threshold and a suppression threshold; soft suppression [27] replaces the hard overlap cut with score decay. The confidence scale itself is not calibrated across conditions [28], which motivates ranking scene statistics against a reference distribution rather than using raw thresholds. In this study all operating-point parameters are selected on validation data only, with the tree-structured Parzen estimator of the Optuna framework [29] used as the search engine.

### D. Content-adaptive inference

Adapting inference to content is an established idea. Dynamic neural networks adapt their depth, width, or resolution to the input [30], and video analytics systems adapt configurations such as resolution and frame rate to the content and to the available resources [31]. These systems either learn the adaptation policy or profile many configurations online.

The controller studied here is deliberately simpler: it is training-free, causal, reads five scalar cues, and changes only three detector parameters, so that its behavior can be audited and its benefit attributed.

## E. Evaluation

Tracking accuracy is reported with the CLEAR multi-object tracking accuracy (MOTA) [32], the identity F1 score (IDF1) [33], and the higher order tracking accuracy (HOTA) [34], which balances detection and association, following the MOTChallenge practice [35]. Uncertainty is quantified by a paired bootstrap over sequences [36].

## III. Problem Formulation

Let  $I_1, I_2, \dots$  be the frames of a video. A detector  $D$  applied to frame  $I_t$  with operating point  $\theta_t = (c_t, \eta_t, r_t)$  returns a set of boxes with confidence scores: the image is resized so that its long side equals  $r_t$  pixels, candidate boxes are produced, boxes overlapping a higher-scoring box by more than  $\eta_t$  are suppressed, and boxes with score below  $c_t$  are discarded. An online tracker  $T$  consumes the kept boxes and outputs identified tracks. A controller  $\pi$  chooses the operating point,

$$\theta_t = \pi(I_{1:t}, \mathcal{D}_{1:t-1}), \quad \theta_t \in \Theta = \mathcal{C} \times \mathcal{H} \times \mathcal{R}, \quad (1)$$


where  $\mathcal{D}_{1:t-1}$  are the detections of earlier frames and  $\mathcal{C}, \mathcal{H}, \mathcal{R}$  are the admissible sets of confidence thresholds, suppression thresholds, and resolutions. Equation (1) makes the controller causal: it may use the current image and past detections but no future frame and no ground truth.

The design objective is to maximize tracking accuracy on held-out sequences subject to a processing-throughput constraint and without increasing identity fragmentation:

$$\max_{\pi} \text{Acc}(\pi) \quad \text{s.t.} \quad \text{FPS}(\pi) \geq 25, \quad \text{IDS}(\pi) \leq \text{IDS}_{\text{ref}}, \quad (2)$$

where FPS is the number of processed frames per second and IDS the number of identity switches. The accuracy term and the reference are made concrete in Sec. VIII: MOTA was the selection objective, HOTA is the primary reporting metric, and  $\text{IDS}_{\text{ref}}$  is the identity-switch count of the hand-designed controller on validation. Because held-out accuracy cannot be observed during design, Eq. (2) is solved on a validation split, and the held-out test measures how well that solution transfers.

Three restrictions define the scope. The controller is training-free: it has no learned weights, only a small number of scalar parameters selected on validation data. The detector and the tracker are fixed and pretrained; neither is fine-tuned on aerial data. The admissible sets are finite and are themselves derived from validation sweeps (Sec. VIII), so that the controller only chooses among operating points with measured validation support.



**Fig. 1 Scene-adaptive operating-point control.** The controller runs every  $K = 10$  frames on a quarter-scale grayscale image and on the detections kept at the previous frame; the chosen operating point is held until the next analysis.

## IV. Methodology

Figure 1 shows the pipeline. Every  $K$  frames, the controller computes five cues from a quarter-scale grayscale copy of the current frame and from the detections of the previous frame, ranks them against a calibration set, combines them into the raw index  $s_j$ , smooths the index over the last  $W$  analyses, and maps the smoothed index  $S_t$  to the operating point used until the next analysis.

### A. Scene cues

Five cues are computed at each analysis. They were chosen because each targets a failure mode of aerial detection and each costs at most a few image passes on a quarter-resolution frame.

*Crowd.* The number  $N_{t-1}$  of detections kept at the previous frame measures density. It is a detector-derived quantity, not a ground-truth count.

*Tiny-object ratio.* The fraction of the previous frame’s kept detections whose box area is below  $32 \times 32$  pixels,

$$q_{\text{tiny}} = \frac{1}{N_{t-1}} \sum_{i=1}^{N_{t-1}} \mathbf{1}[a_i < 32^2], \quad (3)$$

with  $q_{\text{tiny}} = 0$  when  $N_{t-1} = 0$ , measures how much of the scene is at the scale where detectors lose recall [14].

*Edge strength.* The mean Sobel gradient magnitude  $\gamma$  of the quarter-scale grayscale image measures texture and clutter.

*Low light.* The mean gray level  $\beta$  of the same image measures illumination.

*Blur.* The variance  $\nu$  of the Laplacian of the same image is a standard focus measure [37]; low values indicate blur.

## B. Label-free calibration

Raw cue values have incomparable units and dataset-dependent ranges. Each cue is therefore mapped to  $[0, 1]$  by its empirical distribution function on a calibration set,

$$\hat{F}_k(x) = \frac{1}{M} \left| \{m : x_m^{(k)} \leq x\} \right|, \quad (4)$$

where  $x_m^{(k)}$ ,  $m = 1, \dots, M$ , are the calibration values of cue  $k$ . The normalized cues are

$$q_{\text{crowd}} = \hat{F}_N(N_{t-1}), \quad q_{\text{edge}} = \hat{F}_\gamma(\gamma), \quad q_{\text{light}} = 1 - \hat{F}_\beta(\beta), \quad q_{\text{blur}} = 1 - \hat{F}_\nu(\nu), \quad (5)$$

and  $q_{\text{tiny}}$  is already a fraction. The calibration set contains  $M = 288$  analysis samples taken every ten frames from the seven VisDrone validation sequences, with image statistics from the frames and crowd counts from the fixed detector run at a single anchor operating point (960 pixels, confidence 0.35, suppression threshold 0.35; person, car, bus, and truck classes). No ground-truth count enters the calibration. On this set the median gray level is 107.6 (interquartile range 85.4–185.1), the median Laplacian variance is 1980, the median gradient magnitude is 66.2, and the median detection count is 10.5 (maximum 41).

The ranking makes the index invariant to monotone rescaling of each cue, and it makes the weights interpretable: a weight multiplies a quantile, not a unit-dependent value. It also means that the index describes a scene relative to the calibration distribution, so that a scene is "complex" if it is more cluttered, darker, blurrier, denser, or has more tiny objects than typical validation scenes.

## C. Scene complexity index


The raw index of analysis  $j$  is a convex combination of the normalized cues,

$$s_j = \text{clip} \left( \sum_k w_k q_k, 0, 1 \right), \quad w_k \geq 0, \quad \sum_k w_k = 1, \quad (6)$$

and the smoothed index is the mean of the last  $W$  raw values,

$$S_t = \frac{1}{\min(j, W)} \sum_{i=\max(1, j-W+1)}^j s_i, \quad t \in [1 + (j-1)K, jK]. \quad (7)$$

Analyses occur at frames  $t = 1, 1 + K, 1 + 2K, \dots$ ; between analyses the index and the operating point are held. With  $K = 10$  and  $W = 7$ , the index at 30 frames per second reflects roughly the last two seconds and changes at most three times per second. The controller state is reset at the start of each sequence.



**Fig. 2** Composition of the scene complexity index. (a) Normalized cue weights of the frozen quality profile Q and balanced profile B. (b) Importance of each search parameter for validation MOTA in the single-objective search that produced Q, as stored with the search record; the two index thresholds dominate.

## V. Scene Complexity Index Analysis

### A. Frozen weights

Figure 2a shows the normalized weights of the two frozen profiles, the quality profile Q (trial 24 of the single-objective search) and the balanced profile B (trial 22 of the multi-objective search). Both place the largest weight on edge strength (0.434 for Q and 0.508 for B), a second tier on the tiny-object ratio (0.222 and 0.175) and on crowd density (0.129 and 0.165), and small weights on low light and blur (0.054 and 0.161 for Q; 0.121 and 0.032 for B). The weight on edge strength agrees with the intuition that textured urban scenes are the hard ones for aerial detection. The two searches disagree on blur and low light, which suggests that those cues are weakly identified by seven validation sequences.

### B. What the search actually tuned

The parameter importance scores stored with the single-objective search (Fig. 2b) are more revealing than the weights. The two thresholds of the resolution rule account for most of the explained variation in validation MOTA (0.575 and 0.161), followed by the easy-end suppression threshold (0.145); every cue weight has importance below 0.03. In other words, the search mostly tuned *how often the controller uses high resolution*, not how the scene is described. The frozen thresholds of Q ( $\tau_m = 0.135$ ,  $\tau_h = 0.287$ ) are low relative to the typical index, so Q runs at 960 pixels on almost every frame: its mean resolution is 957 pixels on validation and 955 pixels on UAVDT. Its suppression threshold is constant ( $\eta_e = \eta_h = 0.35$ ), and only its confidence threshold moves, between 0.30 and 0.40, with a validation mean of 0.341. The optimized quality controller is therefore close to a static operating point at (960, 0.34, 0.35). This observation is the key to the attribution result of Sec. IX.

The balanced profile behaves differently. Its thresholds ( $\tau_m = 0.293$ ,  $\tau_h = 0.666$ ) place typical scenes in the

intermediate and high levels and simple scenes at 512 pixels, giving mean resolutions of 825 pixels on validation, 869 pixels on test-dev, and 829 pixels on UAVDT. Its confidence threshold is constant at 0.40, and its suppression threshold tightens from 0.60 in simple scenes to 0.35 in complex ones (mean 0.47–0.50). Profile B is thus the one that actually switches, and it trades accuracy for fewer identity switches, fewer false positives, and higher throughput.

### C. Feedback through the crowd cue

The crowd and tiny-object cues are computed from detections that already passed the controller’s own confidence filter, which closes a loop. In Q the confidence threshold rises with the index; a higher index therefore removes low-scoring boxes, lowers the next crowd count, and pulls the index back down. This is negative feedback, which keeps the operating point stable. The hand-designed controller of Sec. VI lowers the confidence threshold as the index rises, which admits more boxes and raises the next crowd count: positive feedback, bounded only by the clipping of its thresholds. The moving mean of Eq. (7) and the ten-frame stride damp both loops. The analysis does not change any result, but it explains why the validation search preferred a controller whose confidence threshold increases with complexity.

## VI. Adaptive Control Strategy

### A. Calibrated control laws


The calibrated controller maps the smoothed index to the operating point by two linear laws and one piecewise-constant law:

$$c_t = c_e + S_t (c_h - c_e), \quad (8)$$

$$\eta_t = \eta_e + S_t (\eta_h - \eta_e), \quad (9)$$

$$r_t = \begin{cases} r_0, & S_t < \tau_m, \\ r_1, & \tau_m \leq S_t < \tau_h, \\ r_2, & S_t \geq \tau_h, \end{cases} \quad (10)$$

with resolution levels  $(r_0, r_1, r_2) = (512, 928, 960)$  pixels. The eleven parameters are the five weights, the four endpoints  $c_e, c_h, \eta_e, \eta_h$ , and the two thresholds  $\tau_m, \tau_h$ . Figure 3 plots the three laws for both profiles, and Table 1 lists all frozen configurations.



**Fig. 3** Adaptation maps: confidence threshold, suppression threshold, and input resolution as functions of the smoothed index for the quality profile Q, the balanced profile B, and the hand-designed controller (scene-label adjustments of the hand-designed controller omitted).

**Table 1** Frozen configurations compared in this study. Confidence and NMS entries give the value at  $S = 0$  and at  $S = 1$  (linear in between); resolution is the long-side input size in pixels. All adaptive controllers analyze every  $K = 10$  frames with a moving mean over  $W = 7$  analyses.

| Configuration                                                                                                                                               | Confidence                             | NMS IoU                               | Resolution                          | Tracker |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------|---------------------------------------|-------------------------------------|---------|
| Static default                                                                                                                                              | 0.25                                   | 0.45                                  | 640                                 | default |
| Hand-designed controller                                                                                                                                    | $0.245 - 0.05S$ , clipped to 0.19–0.28 | $0.49 - 0.05S$ , clipped to 0.40–0.52 | 640/736/832 (rules in Sec. VI)      | tuned   |
| Quality profile Q (trial 24)                                                                                                                                | 0.30→0.40                              | 0.35→0.35                             | 512/928/960 at $S \geq 0.135/0.287$ | tuned   |
| Balanced profile B (trial 22)                                                                                                                               | 0.40→0.40                              | 0.60→0.35                             | 512/928/960 at $S \geq 0.293/0.666$ | tuned   |
| Matched static anchor                                                                                                                                       | 0.35                                   | 0.35                                  | 960                                 | tuned   |
| Tracker profiles (ByteTrack): default = high 0.25, low 0.10, new 0.25, buffer 30, match 0.80; tuned = high 0.18, low 0.04, new 0.20, buffer 45, match 0.86. |                                        |                                       |                                     |         |

## B. Hand-designed controller

The hand-designed controller is the earlier version of the same idea and serves as the adaptive baseline. Its index uses fixed saturating transforms instead of calibrated ranks:

$$s_j = \text{clip} \left( 0.30 \min \left( \frac{N_{t-1}}{30}, 1 \right) + 0.20 \min \left( \frac{e}{0.14}, 1 \right) + 0.30 q_{\text{tiny}} + 0.10 \mathbf{1}[\beta < 80] + 0.05 \mathbf{1}[\nu < 180], 0, 1 \right), \quad (11)$$

where  $e$  is the fraction of Canny edge pixels of the quarter-scale image. It also assigns a scene label (night, blur, tiny, crowded, or clear) from the same cues. Its confidence threshold  $0.245 - 0.05S_t$  decreases with complexity, is lowered by a further 0.012 in crowded, tiny, or night scenes, and is clipped to  $[0.19, 0.28]$ ; its suppression threshold  $0.49 - 0.05S_t$  is lowered by 0.012 in blurred scenes and clipped to  $[0.40, 0.52]$ . Its resolution is 832 pixels when  $S_t > 0.60$  or the tiny-object ratio exceeds 0.5, 736 pixels when  $S_t > 0.35$  or the scene is labelled crowded or tiny, and 640 pixels otherwise. It uses the same analysis stride and smoothing window as the calibrated controller ( $K = 10$ ,  $W = 7$ ).

The two controllers encode opposite hypotheses about confidence. The hand-designed controller is recall-oriented:

in hard scenes it admits weaker boxes. The calibrated profiles are precision-oriented: Q raises the confidence threshold in hard scenes, and B holds it at 0.40 while tightening suppression. The validation evidence of Sec. X favors the precision-oriented choice for this detector and tracker.

### C. Update schedule and cost

All controllers analyze the first frame of each sequence and every tenth frame thereafter; the image cues require one resize, one color conversion, two Sobel filters, and one Laplacian on a quarter-scale image, and the detection cues reuse boxes that the tracker already holds. In the matched static run on test-dev, where the controller code path ran without cue analysis, the per-frame controller time was 0.005 ms against 19.5 ms for the detector and 3.6 ms for the tracker. The cue analysis cost of the adaptive runs was not profiled separately; it is included in their end-to-end processing throughput.

## VII. Detection and Tracking Pipeline

### A. Detector

The detector is the Ultralytics YOLOv8n model pretrained on the COCO dataset (Ultralytics library version 8.3.200), run in half precision through PyTorch on one NVIDIA Tesla T4 graphics processor with at most 1000 detections per frame. Only the person, car, bus, and truck classes are kept on VisDrone; car, bus, and truck are kept on UAVDT, whose annotations cover vehicles only. The detector is not fine-tuned on aerial data, which keeps the controller evaluation independent of any training choice but also bounds the absolute accuracy (Sec. XIV).

### B. Tracker

The tracker is the Ultralytics implementation of ByteTrack [3] with score fusion and a nominal frame rate of 30. Two profiles are used (Table 1): the library-style default profile (high threshold 0.25, low threshold 0.10, new-track threshold 0.25, buffer 30 frames, matching threshold 0.80) and a tuned profile selected in earlier validation work (0.18, 0.04, 0.20, 45 frames, 0.86). The tuned profile is part of every adaptive system and of the matched static anchor; the static default uses the default profile.

One interaction between detector and tracker needs to be stated explicitly. Detections below the controller’s confidence threshold are removed before the tracker. Because every confidence threshold used in this study (0.19 or higher) is at or above the tracker’s high threshold, every detection reaching the tracker counts as high-confidence, and all of them except boxes of the hand-designed controller scored between 0.19 and 0.20 are eligible to start a new track. ByteTrack’s second association stage, which recovers occluded objects from low-confidence boxes, therefore receives no input in these configurations. The tuned profile acts mainly through its longer buffer and its looser matching threshold.

**Table 2 Data splits and their role. No test sequence was used for any selection, calibration, or tuning decision.**

| Dataset          | Split      | Sequences | Frames | Role                                             |
|------------------|------------|-----------|--------|--------------------------------------------------|
| VisDrone2019-MOT | validation | 7         | 2,846  | cue calibration, sweeps, ablations, optimization |
| VisDrone2019-MOT | test-dev   | 17        | 6,635  | locked held-out comparison                       |
| UAVDT            | test       | 20        | 16,592 | zero-tuning cross-dataset test                   |

### C. Timing protocol

Throughput is reported as processing frames per second: the number of frames divided by the summed per-frame time of cue analysis, detection (with device synchronization), and tracking. Frames are decoded from JPEG files before timing, and all resolutions a system can request are warmed up before the first timed frame. Decoding is excluded because an onboard pipeline receives decoded frames from the camera. In the matched static run, decoding cost 22.5 ms per frame on average, and throughput including decoding was 21.90 frames per second, below the 25 frames-per-second target. The throughput figures of this paper are therefore processing throughput on already-decoded frames, not end-to-end real-time performance on board. The gate of 25 frames per second was applied to the whole-split processing throughput.

## VIII. Experimental Setup

### A. Data

Table 2 lists the data. All design decisions used the seven sequences of the VisDrone2019 multi-object tracking validation split [4]. The seventeen sequences of the test-dev split were held out until a locked three-system comparison. The twenty sequences of the UAVDT test split [5] were used once, as a zero-tuning cross-dataset test.

### B. Evaluation protocol

Evaluation is class-agnostic. On VisDrone, ground-truth boxes of the pedestrian, car, van, truck, and bus categories are kept when they are marked as considered, with occlusion and truncation levels below 2; all detector classes are pooled, so a van can be matched by any detector box, including one labelled car or truck, since the COCO detector has no van class. This is a research protocol, not the official VisDrone evaluation, so the absolute numbers are not comparable with the VisDrone leaderboard. On UAVDT, a frozen adapter (checked by a self-test on a training sequence, which it reproduced with perfect scores) removes ground truth marked out of scope, normalizes identities, removes predictions strictly contained in ignore regions for the eighteen sequences that provide them, and matches at an overlap of 0.5. All metrics were computed with the TrackEval library pinned to commit 12c8791.

HOTA is the primary metric because it weighs detection and association equally [34]; MOTA, IDF1, identity switches, false negatives, and false positives are reported alongside. Uncertainty is quantified by a paired bootstrap over sequences [36]: sequences are resampled with replacement 5,000 times (seed 42), the metric difference between two

systems is recomputed on each resample from the pooled counts, and the 2.5th and 97.5th percentiles form the 95% interval. A win, tie, or loss count per sequence is reported as a complementary, distribution-free summary.

### C. Selection chain

The selection chain used the validation split only.

*Operating-point sweeps.* With the tuned tracker, three one-at-a-time sweeps measured validation accuracy and throughput for resolutions from 512 to 960 pixels in steps of 32, confidence thresholds from 0.05 to 0.50, and suppression thresholds from 0.30 to 0.80. Confidence and suppression values were kept if they met the throughput gate and at least the reference accuracy of the sweep, which retained confidence values 0.25–0.45 and suppression values 0.30–0.70. Three resolution levels were retained by a predeclared rule: a low-compute feasible level, the highest-quality feasible level, and the highest-quality distinct intermediate level, which gave 512, 960, and 928 pixels.

*Temporal settings.* A grid of smoothing windows  $W \in \{1, 3, 5, 7, 9\}$  and analysis strides  $K \in \{1, 5, 10, 15, 20\}$  was evaluated with the hand-designed controller, and  $W = 7, K = 10$  was frozen.

*Quality profile.* A tree-structured Parzen estimator search of at most 50 trials explored the five raw weights (normalized to sum to one), the confidence and suppression endpoints within the retained sets, and the two resolution thresholds. The selection rule, fixed before the search, required the throughput gate and no more identity switches than the hand-designed controller on validation (271), and then selected the highest MOTA, breaking ties by fewer identity switches, then HOTA, IDF1, and throughput. Trial 24 was selected, with validation HOTA 36.11%, MOTA 23.04%, IDF1 40.76%, 270 identity switches, and 37.17 frames per second.

*Balanced profile.* A second, multi-objective search (seed 42) maximized MOTA and minimized identity switches with only the throughput gate. Of 50 planned trials, 49 completed. The Pareto front ranged from trial 16 (MOTA 22.83%, 308 switches) to trial 8 (11.54%, 114 switches). A balanced score, declared before the search, averaged min–max normalized MOTA and normalized identity-switch reduction, and it selected trial 22 (validation MOTA 19.33%, HOTA 31.65%, 168 switches, 51.41 frames per second).

### D. Held-out protocol

Before the test split was opened, a protocol file fixed the three compared systems (static default, hand-designed controller, quality profile Q), their configuration hashes, the calibration hash, the ground-truth filter, the evaluator commit, the throughput gate, and the requirement of a T4 processor. Each system ran as a separate worker in its own session and wrote an immutable result file; the merge step refused to run until all three existed. The quality-profile worker resumed from per-sequence checkpoints after an interruption, with no parameter change. The balanced profile was frozen after its selection and tested with the same machinery; its first test attempt was interrupted before any result was saved, and an exact technical rerun of the frozen configuration produced the reported result. The UAVDT test was

**Table 3 Held-out results (percent for HOTA, MOTA, IDF1; counts for IDS, FN, FP; processing frames per second on one NVIDIA Tesla T4, each system measured in its own session). No entry is marked as best because the profiles target different trade-offs.**

| System                                                              | HOTA  | MOTA  | IDF1  | IDS  | FN     | FP    | FPS                |
|---------------------------------------------------------------------|-------|-------|-------|------|--------|-------|--------------------|
| <i>VisDrone2019-MOT test-dev (17 sequences, 6,635 frames)</i>       |       |       |       |      |        |       |                    |
| Static default                                                      | 28.43 | 19.73 | 32.72 | 1235 | 155051 | 16308 | 36.53              |
| Hand-designed controller                                            | 32.70 | 23.24 | 39.52 | 1061 | 138967 | 25025 | 41.96              |
| Quality profile Q                                                   | 33.84 | 26.95 | 41.55 | 1184 | 136858 | 19029 | 38.98              |
| Balanced profile B                                                  | 31.22 | 23.79 | 37.87 | 919  | 149307 | 13632 | 46.02              |
| Matched static anchor                                               | 33.93 | 27.00 | 41.77 | 1197 | 136391 | 19374 | 43.23 <sup>a</sup> |
| <i>UAVDT test (20 sequences, 16,592 frames), no tuning on UAVDT</i> |       |       |       |      |        |       |                    |
| Static default                                                      | 24.08 | 13.84 | 27.89 | 558  | 270189 | 22974 | 65.01              |
| Quality profile Q                                                   | 28.39 | 17.40 | 34.45 | 321  | 258376 | 22896 | 58.35              |
| Balanced profile B                                                  | 26.93 | 16.12 | 32.01 | 308  | 266894 | 18758 | 61.46              |

<sup>a</sup>Run after the freeze for attribution only; frames per second measured in a separate T4 session.

**Table 4 Paired sequence-level bootstrap (5,000 resamples, seed 42): observed difference and 95% percentile interval. HOTA, MOTA and IDF1 in percentage points; IDS reduction in switches (positive = fewer switches for the first system).**

| Data     | Comparison                 | $\Delta$ HOTA        | $\Delta$ MOTA        | $\Delta$ IDF1         | IDS reduction     |
|----------|----------------------------|----------------------|----------------------|-----------------------|-------------------|
| VisDrone | Q vs static default        | +5.41 [+4.13, +6.90] | +7.22 [+5.44, +9.29] | +8.82 [+6.90, +11.05] | +51 [−114, +219]  |
| VisDrone | Q vs hand-designed         | +1.14 [+0.14, +2.27] | +3.71 [+2.39, +5.19] | +2.03 [+0.59, +3.77]  | −123 [−244, −14]  |
| VisDrone | Q vs matched static anchor | −0.09 [−0.25, +0.06] | −0.05 [−0.25, +0.16] | −0.23 [−0.50, +0.02]  | +13 [−23, +46]    |
| VisDrone | B vs static default        | +2.79 [+0.75, +4.66] | +4.06 [+0.95, +6.89] | +5.15 [+2.03, +8.08]  | +316 [+175, +470] |
| VisDrone | B vs Q                     | −2.62 [−4.00, −1.69] | −3.16 [−6.14, −1.35] | −3.68 [−5.88, −2.24]  | +265 [+82, +502]  |
| UAVDT    | Q vs static default        | +4.30 [+2.74, +5.71] | +3.56 [+1.45, +5.81] | +6.57 [+3.81, +8.82]  | +237 [+104, +377] |
| UAVDT    | B vs static default        | +2.85 [+1.24, +4.35] | +2.28 [+0.72, +4.16] | +4.13 [+1.63, +6.49]  | +250 [+110, +404] |
| UAVDT    | B vs Q                     | −1.46 [−2.19, −0.91] | −1.28 [−2.44, −0.20] | −2.44 [−3.74, −1.37]  | +13 [−20, +50]    |

run for the static default and both frozen profiles under a separate lock that forbade any tuning, recalibration, or selection on UAVDT.

After these tests were complete, the matched static anchor (960 pixels, confidence 0.35, suppression 0.35, tuned tracker, no adaptation) was evaluated on test-dev and paired with the quality profile. It was not part of the preregistered comparison: it serves only to attribute the measured gain, and its throughput was measured in a separate session.

## IX. Main Results

### A. Held-out test-dev

Table 3 gives the held-out results and Table 4 the paired bootstrap intervals.

Against the static default, the quality profile raised HOTA from 28.43% to 33.84% ( $\Delta = +5.41$ , 95% interval [+4.13, +6.90]), MOTA by 7.22 points, and IDF1 by 8.82 points, and it won on all 17 sequences. The gain comes mostly from recall: false negatives fell by 18,193, while false positives rose by 2,721. Identity switches fell by 51, but that interval includes zero. Throughput was 38.98 frames per second against 36.53 for the default, both above the gate.

Against the hand-designed controller, the quality profile gained 1.14 HOTA points ( $[+0.14, +2.27]$ ), 3.71 MOTA points, and 2.03 IDF1 points, and it won on 13 of 17 sequences. It produced 5,996 fewer false positives and 2,109 fewer false negatives, but 123 more identity switches, and that difference is significant ( $[-244, -14]$  in switch reduction). The hand-designed controller therefore remains the identity-switch leader of the preregistered comparison, and the quality profile is not better on every metric.

The balanced profile sits at a different point of the trade-off. On test-dev it had the fewest identity switches (919), the fewest false positives (13,632), and the highest throughput (46.02 frames per second), with HOTA 31.22%. It was significantly better than the static default on every metric, including 316 fewer identity switches ( $[+175, +470]$ ), and significantly worse than the quality profile in HOTA ( $-2.62$ ), MOTA, and IDF1, while having 265 fewer switches ( $[+82, +502]$ ).

### **B. Cross-dataset test without tuning**

On UAVDT, no parameter was changed. The quality profile raised HOTA from 24.08% to 28.39% ( $[+2.74, +5.71]$ ), MOTA by 3.56 points, and IDF1 by 6.57 points, and it reduced identity switches from 558 to 321 ( $[+104, +377]$ ), winning 15 of 20 sequences with 2 ties and 3 losses. The balanced profile improved over the default by 2.85 HOTA points ( $[+1.24, +4.35]$ ) with 250 fewer switches, but it was significantly below the quality profile in HOTA ( $-1.46$ ), MOTA, and IDF1, and its 13-switch advantage over the quality profile is not significant ( $[-20, +50]$ ). Both profiles ran slower than the static default on UAVDT (58.35 and 61.46 against 65.01 frames per second), because both use higher mean resolutions than the default’s 640 pixels.

Across both datasets, the quality profile is the accuracy leader of the frozen systems, and the balanced profile is a reproducible identity, false-positive, and throughput trade-off on VisDrone whose identity advantage over the quality profile does not transfer significantly to UAVDT.

### **C. Attribution: the matched static anchor**

The quality profile runs at the top resolution on almost every frame with a confidence threshold near 0.34 and a constant suppression threshold of 0.35 (Sec. V). Its natural control is therefore the static operating point (960, 0.35, 0.35) with the same tuned tracker, which is also the anchor of the cue calibration. On validation, that anchor reached 36.08% HOTA against 36.11% for the full quality profile (Table 5). On test-dev, the anchor reached 33.93% HOTA against 33.84% for the quality profile:  $\Delta = -0.09$  for the profile, with interval  $[-0.25, +0.06]$  (Table 4). MOTA, IDF1, and identity switches likewise differ by amounts whose intervals include zero.

On validation, the two parts of that configuration can be separated with the ablation runs of Table 5: at the default operating point, the tuned tracker profile added 1.03 HOTA points (H0 to H1), and moving the detector to the anchor operating point with the same tracker added a further 5.21 points (H1 to C0). Most of the calibration gain therefore

comes from the detector operating point, and a smaller part from the tracker profile; this decomposition is a validation result, because the tracker profiles were not separated on the held-out splits.

The held-out gains of the quality profile over the static default (+5.41 HOTA) and over the hand-designed controller (+1.14) are therefore attributable to the validation-calibrated operating point, namely high resolution, a precision-oriented confidence threshold, a tight suppression threshold, and the tuned tracker profile, and not to frame-to-frame switching. The frozen protocol and the post hoc anchor jointly support that conclusion; the anchor was not available as a preregistered competitor, which is why the paper reports it as an attribution analysis rather than as a fourth system of the locked comparison.

#### **D. Validation-to-test transfer**

All selection used seven validation sequences, and the controller has eleven parameters, so overfitting to the validation split is a natural concern. Three observations bound it. First, the gain of the quality profile over the static default was 6.27 HOTA points on validation (36.11 against 29.84 percent, Table 5), 5.41 points on the seventeen held-out test-dev sequences, and 4.30 points on UAVDT without any retuning, so most of the validation gain transferred to unseen sequences and to a second dataset. Second, the search concentrated its effect on the two resolution thresholds (Sec. V), and the selected controller behaves like a static operating point with three parameters, which is the configuration that the matched static anchor reproduces. Third, the one constraint that did not transfer, the identity-switch bound, is reported as a failure in Sec. XII. The transfer supports the calibrated operating point; it does not show that the eleven parameters are individually identified, and the weakly identified cue weights (Sec. V) should not be read as a general ranking of scene cues.

#### **E. Per-sequence behavior**


Figure 4 shows the per-sequence HOTA differences. On test-dev, the gain of the quality profile over the default is positive everywhere and ranges from +1.07 to +13.10 points; the largest gain occurs on the sequence where all systems are weakest (default HOTA 7.10%). Against the hand-designed controller, four sequences are losses, the largest by 4.69 points. On UAVDT, the gains of both profiles are positive on most sequences, with a tail of losses discussed in Sec. XII.

## **X. Ablation**

Table 5 and Fig. 5 decompose both controllers on the validation split.

#### **A. Hand-designed controller**

Starting from the static default (H0, 29.84% HOTA), tuning the tracker (H1) added 1.03 HOTA points and removed 66 identity switches at unchanged throughput. Adaptive confidence and suppression at a fixed 640 pixels (H2) added a further 0.52 points. Adaptive resolution alone, with static confidence and suppression (H2R), reached 32.45%, 1.59



**Fig. 4** Per-sequence HOTA differences, sorted. (a) VisDrone test-dev: quality profile minus static default and minus the hand-designed controller. (b) UAVDT: quality and balanced profiles minus static default, without any tuning on UAVDT.


**Table 5** Component ablations on VisDrone2019-MOT validation (7 sequences). FPS is processing throughput on a T4; mean resolution, confidence and NMS are averaged over frames.

| Stage                                           | Components                                 | HOTA  | MOTA  | IDF1  | IDS | FPS  | Res. | Conf. | NMS   |
|-------------------------------------------------|--------------------------------------------|-------|-------|-------|-----|------|------|-------|-------|
| <i>Hand-designed controller</i>                 |                                            |       |       |       |     |      |      |       |       |
| H0                                              | static default, default tracker            | 29.84 | 17.63 | 30.89 | 283 | 44.1 | 640  | 0.250 | 0.450 |
| H1                                              | + tuned tracker                            | 30.86 | 17.80 | 32.85 | 217 | 44.5 | 640  | 0.250 | 0.450 |
| H2                                              | + adaptive conf./NMS (640 px)              | 31.38 | 17.64 | 33.73 | 210 | 40.7 | 640  | 0.218 | 0.474 |
| H2R                                             | tuned tracker + adaptive resolution only   | 32.45 | 18.42 | 35.67 | 241 | 39.6 | 722  | 0.250 | 0.450 |
| H3                                              | full hand-designed controller              | 33.06 | 18.16 | 36.30 | 271 | 38.0 | 725  | 0.218 | 0.474 |
| <i>Optimized controller (quality profile Q)</i> |                                            |       |       |       |     |      |      |       |       |
| C0                                              | static anchor 960/0.35/0.35, tuned tracker | 36.08 | 23.08 | 40.60 | 266 | 37.9 | 960  | 0.350 | 0.350 |
| C1                                              | + adaptive confidence                      | 36.11 | 23.01 | 40.73 | 271 | 35.3 | 960  | 0.341 | 0.350 |
| C2                                              | + adaptive NMS                             | 36.11 | 23.01 | 40.73 | 271 | 35.5 | 960  | 0.341 | 0.350 |
| C3                                              | + adaptive resolution (full Q)             | 36.11 | 23.04 | 40.76 | 270 | 35.4 | 957  | 0.341 | 0.350 |

points above H1, at a mean resolution of 722 pixels. The full controller (H3) reached 33.06%, 2.20 points above H1, at 38.0 frames per second against 44.5 for H1. For the hand-designed controller, therefore, scene-driven switching does add accuracy, and most of it comes from resolution. The operating range of that controller is narrow and low (640–832 pixels), and its adaptive resolution moves the operating point into the range where accuracy grows fastest with resolution (Table 7).

## B. Calibrated controller

For the quality profile the picture is different. The static anchor (C0) already reached 36.08% HOTA, 3.01 points above the full hand-designed controller. Adding adaptive confidence (C1) changed HOTA by +0.03 points and identity switches by +5. Adding adaptive suppression (C2) changed nothing, as expected, because the profile’s two suppression endpoints are equal; the identical metrics confirm that the ablation machinery isolates components exactly. Adding adaptive resolution (C3) changed HOTA by less than 0.01 points, because the index almost never falls below the



**Fig. 5 Validation ablations.** (a) Hand-designed controller: H0 static default with default tracker; H1 tuned tracker; H2 adaptive confidence and suppression at 640 pixels; H2R adaptive resolution only; H3 full controller. (b) Calibrated controller: C0 static anchor; C1 adaptive confidence; C2 adaptive confidence and suppression; C3 full quality profile; H3 shown for reference. (c) Temporal grid of the hand-designed controller, selected setting outlined.

**Table 6 Temporal settings of the hand-designed controller on VisDrone2019-MOT validation: HOTA (%) / processing FPS for smoothing window  $W$  (analyses) and analysis stride  $K$  (frames). The selected setting is  $W = 7$ ,  $K = 10$ .**

| $W \setminus K$ | 1            | 5            | 10                  | 15           | 20           |
|-----------------|--------------|--------------|---------------------|--------------|--------------|
| 1               | 32.98 / 34.3 | 33.07 / 45.3 | 33.02 / 45.9        | 33.00 / 47.3 | 33.05 / 47.9 |
| 3               | 32.98 / 37.0 | 33.07 / 46.8 | 32.98 / 47.8        | 32.99 / 48.9 | 32.98 / 48.1 |
| 5               | 33.01 / 36.1 | 33.07 / 45.7 | 33.02 / 48.1        | 32.98 / 48.7 | 32.82 / 49.0 |
| 7               | 32.99 / 36.0 | 33.06 / 45.8 | <b>33.06 / 43.9</b> | 33.00 / 48.5 | 32.74 / 49.8 |
| 9               | 32.89 / 37.3 | 33.12 / 47.8 | 33.02 / 48.3        | 32.93 / 43.5 | 32.82 / 45.1 |


resolution thresholds. The components of the calibrated controller are therefore measurable but negligible, which the held-out anchor comparison of Sec. IX confirms on unseen data.

### C. Temporal settings

Across the 25 temporal settings, HOTA of the hand-designed controller varied only between 32.74% and 33.12% (Table 6); the frozen setting  $W = 7$ ,  $K = 10$  gave 33.06%. Analyzing every frame ( $K = 1$ ) cost 6 to 14 frames per second relative to  $K \geq 5$  in the same grid without an accuracy benefit, so infrequent analysis is the right design for this controller. Temporal parameters are second-order for accuracy and first-order for cost.

## XI. Accuracy–Runtime Analysis

Figure 6 and Table 7 relate accuracy and throughput.



**Fig. 6 Accuracy versus processing throughput on one Tesla T4. (a) Validation operating points: resolution sweep at confidence 0.25 and confidence sweep at 960 pixels. (b) Held-out systems on VisDrone test-dev (circles) and UAVDT (triangles); the matched static anchor (diamond) was timed in a separate session. The dotted line marks 25 frames per second.**

**Table 7 Accuracy–throughput trade-off of the detector input resolution on VisDrone2019-MOT validation (confidence 0.25, NMS IoU 0.70, tuned tracker; T4 processing throughput and 95th-percentile per-frame latency).**

| Resolution | HOTA  | MOTA  | IDF1  | IDS | FPS  | P95 latency (ms) |
|------------|-------|-------|-------|-----|------|------------------|
| 512        | 26.05 | 13.93 | 26.00 | 216 | 62.7 | 26.1             |
| 576        | 27.77 | 15.67 | 29.19 | 240 | 57.6 | 29.0             |
| 640        | 30.43 | 16.96 | 32.23 | 326 | 53.5 | 32.1             |
| 704        | 31.56 | 17.35 | 34.26 | 322 | 49.8 | 33.6             |
| 768        | 32.09 | 18.15 | 34.45 | 359 | 47.3 | 36.9             |
| 832        | 33.53 | 18.92 | 37.02 | 413 | 45.3 | 37.5             |
| 896        | 35.07 | 20.17 | 39.08 | 399 | 45.0 | 38.2             |
| 928        | 35.10 | 20.67 | 39.70 | 416 | 42.9 | 38.6             |
| 960        | 36.03 | 20.87 | 40.25 | 420 | 42.0 | 40.9             |

### A. Resolution dominates both accuracy and cost

In the validation sweep, raising the input resolution from 512 to 960 pixels raised HOTA from 26.05% to 36.03% and lowered throughput from 62.7 to 42.0 frames per second, with the 95th-percentile per-frame latency growing from 26.1 to 40.9 ms. The confidence threshold mostly moves the balance between false positives and false negatives: at 960 pixels, HOTA stayed within 36.9% and 35.6% for thresholds between 0.05 and 0.35 and fell to 30.6% at 0.50, while false positives fell from 18,056 to 3,426. The very low thresholds cost throughput through the tracker (24.8 frames per second at 0.05), which is why they failed the gate. The suppression threshold had almost no effect between 0.30 and 0.55 and hurt identity switches above 0.60 (266 switches at 0.30–0.50, 467 at 0.80). The calibrated profiles lie close to the best sweep points, which is consistent with the finding that their benefit is the operating point itself.

### **B. Held-out throughput**

All held-out systems met the gate on the whole split, and every individual test-dev sequence was processed above 25 frames per second by every system (the slowest sequence of the quality profile ran at 27.5). The balanced profile was the fastest adaptive system (46.02 frames per second on test-dev), because it spends part of the time at 512 pixels.

### **C. Measurement variability**

Throughput measured in separate sessions on the same processor class varies more than the differences between some systems. The static anchor configuration was timed at 37.92 frames per second in the component ablation and at 44.78 in the suppression sweep, with identical tracking metrics. The full hand-designed controller was timed at 37.96, 42.09, and 43.95 frames per second in three validation runs of the same configuration. Throughput differences of a few frames per second between systems measured in different sessions, including the 2.97 frames-per-second difference between the quality profile and the hand-designed controller on test-dev, should therefore not be interpreted as algorithmic differences. The accuracy metrics are unaffected by this variability.

## **XII. Failure Analysis**

### **A. Identity switches**

The largest failure of the quality profile is identity fragmentation relative to the hand-designed controller on test-dev (123 more switches, significant). On validation the two were matched by construction (270 against 271 switches), so the constraint of Eq. (2) did not transfer. A plausible mechanism is the tracker interaction of Sec. VII: with a precision-oriented confidence threshold, boxes that briefly drop below the threshold disappear from the tracker entirely, and ByteTrack’s low-score recovery cannot bridge them. The balanced profile, which optimized identity switches directly, did control them on both datasets.

### **B. Sequences where adaptation loses**

Against the hand-designed controller, the quality profile lost on four of the 17 test-dev sequences by 0.10, 0.79, 1.87, and 4.69 HOTA points. On UAVDT, the quality profile lost to the static default on three of 20 sequences, by 1.48 to 2.29 points. The balanced profile’s worst UAVDT sequence lost 12.52 points against the default, with false negatives rising from 3,477 to 4,256: its fixed confidence threshold of 0.40 removes true detections that the default’s 0.25 keeps. A fixed precision-oriented threshold is therefore a risk under a domain shift that lowers detector confidence.

### **C. Scenes outside the detector’s reach**

Three UAVDT sequences had HOTA below 0.5% for every system, and in one of them every system missed all 75,412 ground-truth boxes. No operating point recovers objects that the detector does not produce at any threshold.

These sequences were kept in all aggregates. Their cause (detector coverage, class definitions, or annotation conventions) was not diagnosed within the frozen protocol, so the paper does not attribute them.

#### **D. Why switching did not help**

The calibrated search removed the conditions under which switching helps. The hand-designed controller switches between resolutions on the steep part of the accuracy–resolution curve (640–832 pixels), where a scene-dependent increase buys real recall. The calibrated profile starts from the top of that curve; above 928 pixels, the sweep gains only 0.92 HOTA points up to 960 pixels, so downward switches cost accuracy and upward switches are impossible. With seven validation sequences and a throughput gate that 960 pixels already meets, the search’s best strategy was to stay high. Scene-driven switching should therefore be expected to pay off when the throughput budget does not admit the best static operating point, which is the regime of the hand-designed controller and of lower-power processors.

#### **E. Do the scene cues predict where switching helps?**

A separate audit, run later in the same research program on the same seven validation sequences, tested the scene cues directly. It used a different, score-gated controller with resolutions of 640, 736, and 832 pixels, two detectors (YOLOv8n and RT-DETR-L [38]), and ByteTrack, so its numbers are not comparable with the tables of this paper. First, the resolution rule of the hand-designed index, with added hysteresis, used a pixel budget between those of fixed 640 and fixed 736 pixels and reached an accuracy between theirs on both detectors, below fixed 736; no mixture of fixed resolutions was run, so a mixture of equal pixel cost was not measured directly. Second, against the frame-wise benefit of 832 over 640 pixels, computed offline from the ground truth, no causal cue (detection count, tiny-object ratio, box area, score gap, edge strength, brightness, or blur) reached a within-sequence Spearman correlation of 0.3 in magnitude, and the signs changed between sequences and detectors. Third, allocating the higher resolution by any single cue or by the index did not beat random allocation to half of the analysis steps on both detectors. The cue rules did not match the pixel budget of random allocation exactly, and the only rules that exceeded it on one detector also used more pixels: the tiny-object ratio (1.26 HOTA points higher on YOLOv8n with a 19% larger budget, 0.45 points lower on RT-DETR-L) and, by 0.04 points on YOLOv8n, the hand-designed index. A cue-utility analysis with leave-one-sequence-out cross-validation and a permutation null found signal above chance for both detectors only in cues describing the detector output, the tracker state, and image motion, and none in the image cues (edge strength, brightness, blur), although single training folds occasionally selected an image cue; controllers learned from the selected cues still scored lower than fixed global parameters on held-out sequences (significantly only for YOLOv8n in both attempts).

The audit also qualifies the hand-designed ablation of Sec. X: that controller’s adaptive resolution spent more pixels on average (722 against 640), and in the audit the scene-driven rule performed within the range of fixed resolutions of

similar pixel cost, which suggests, for a different controller, that the gain came mainly from the larger pixel budget. The audit thus agrees with the attribution result: in this data, the scene cues do not identify the frames in which a different operating point pays off.

### XIII. Cross-Pipeline Consistency

Two experiments run separately in the same research program, with other detectors, trackers, and policies, test whether the main finding recurs (Table 8). They are supplementary: neither uses Q unchanged and their ground-truth filters differ from Sec. VIII, so only paired differences are compared. Both were scored after the fact from archived tracker outputs, without rerunning any detector or tracker: the archived motmetrics results (false positives, false negatives, identity switches, MOTA, IDF1) were first reproduced exactly, and HOTA was then computed with TrackEval (commit 12c8791) and compared by a paired sequence bootstrap (10,000 resamples, seed 0). HOTA values come from TrackEval and MOTA and IDF1 values from motmetrics, never mixed within a comparison.

*U2MOT on VisDrone test-dev (held-out).* The U2MOT tracker [39] with its authors' YOLOX-X detector [40] was run on the seventeen test-dev sequences. The comparator is the configuration released by the U2MOT authors, the author-calibrated static operating point; it is not a matched static anchor in the sense of Sec. IX. The controller is a separately recalibrated member of the same design family, not Q transferred unchanged (Table 8, note b); it was frozen, with hashes, before the test-dev run. Evaluation follows the protocol of the U2MOT repository, not the official VisDrone toolkit. In a post hoc TrackEval HOTA rescore of the frozen outputs, the controller reached 54.94% against 55.00% ( $\Delta = -0.06$ ,  $[-0.31, +0.22]$ ; 9 wins, no ties, 8 losses; leave-one-out range  $-0.14$  to  $+0.04$ ). With motmetrics, MOTA was 53.77% against 53.87% and IDF1 69.85% against 69.78%. At this sample size, HOTA did not change measurably.

*SparseTrack on MOT17 (in-sample validation).* SparseTrack [41] with a YOLOX detector was run on the validation half of the seven MOT17 training sequences [42]. Its rule switches the suppression threshold between 0.70 and 0.80 with a smoothed edge cue (note c); the cue threshold was selected by a leave-one-sequence-out search on the same seven sequences, and a static threshold of 0.75 was named as the comparator in the rule's freeze record, whose primary metric was MOTA. Against it, HOTA was 69.10% against 68.96% ( $\Delta = +0.13$ ,  $[+0.02, +0.24]$ ; 6 wins, no ties, 1 loss; leave-one-out range  $+0.05$  to  $+0.18$ ); with motmetrics, MOTA was 76.93% against 76.89% and IDF1 81.53% against 81.49%. The default threshold of 0.70 in the tracker's configuration, run in a different batch, reached 69.17%, the highest HOTA of the three static settings rescored, and did not differ measurably from the rule ( $\Delta = -0.07$ ,  $[-0.54, +0.19]$ ; 4 wins, no ties, 3 losses). The frozen rule has no held-out score.

In the primary study and in the U2MOT experiment, scene switching therefore did not measurably outperform a well-calibrated static operating point. The SparseTrack gain suggests that switching may help relative to some static settings, but it was measured in-sample and did not exceed the tracker's default static setting in HOTA; it does not establish held-out generalization. The pipelines share an outcome, not one transferred controller. No throughput is

**Table 8 Cross-pipeline consistency: scene switching versus a static operating point.  $\Delta$ HOTA is adaptive minus static (TrackEval, percentage points) with its 95% paired sequence bootstrap interval (row 1: 5,000 resamples, seed 42; rows 2–3: 10,000 resamples, seed 0); W/T/L counts the sequences won, tied, and lost by the adaptive system. Detectors, trackers, ground-truth filters, and policies differ between rows, so absolute accuracies are not comparable across rows.**

| Pipeline                         | Dataset/split                       | Adaptive mechanism                                                              | Static comparator                                                            | $\Delta$ HOTA | 95% CI         | W/T/L            | Evidence status                                                                                | Interpretation                                                                                      |
|----------------------------------|-------------------------------------|---------------------------------------------------------------------------------|------------------------------------------------------------------------------|---------------|----------------|------------------|------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|
| YOLOv8n + ByteTrack (this study) | VisDrone2019 test-dev, 17 sequences | Quality profile Q: index $\rightarrow$ confidence, NMS, resolution              | Matched static anchor (conf. 0.35, NMS 0.35, 960 px, tuned tracker)          | −0.09         | [−0.25, +0.06] | n/a <sup>a</sup> | Held-out; Q frozen before test; anchor run post hoc                                            | No measurable change                                                                                |
| YOLOX-X + U2MOT                  | VisDrone2019 test-dev, 17 sequences | Recalibrated index tiers $\rightarrow$ confidence, NMS, input size <sup>b</sup> | Author-calibrated static operating point (conf. 0.09, NMS 0.70, 1600×896 px) | −0.06         | [−0.31, +0.22] | 9/0/8            | Held-out; controller frozen before test; post hoc TrackEval HOTA rescore of the frozen outputs | No measurable change                                                                                |
| YOLOX + SparseTrack              | MOT17 validation half, 7 sequences  | Edge cue $\rightarrow$ NMS 0.70 or 0.80 <sup>c</sup>                            | Static NMS 0.75 (named in the freeze record)                                 | +0.13         | [+0.02, +0.24] | 6/0/1            | In-sample validation; threshold selected on the same 7 sequences                               | Small in-sample gain over NMS 0.75, none over default NMS 0.70 <sup>c</sup> ; not held-out evidence |

<sup>a</sup>Per-sequence values of the matched static anchor were not archived.

<sup>b</sup>U2MOT authors’ YOLOX-X checkpoint; tracker thresholds 0.5/0.1, matching 0.8, buffer 15 in both runs. Index: cue weights of Q, cue transforms of Eq. (11); boundaries 0.519 and 0.590 (label-free terciles on VisDrone validation). Operating points (confidence, NMS, input): (0.15, 0.70, 1280×704), (0.12, 0.65, 1440×800), (0.09, 0.60, 1600×896), selection not recorded. The comparator is the separate configuration released by the U2MOT authors, not the third operating point: it shares that point’s confidence and input size but uses NMS 0.70, not 0.60. Evaluation: pedestrian, car, van, truck, bus merged; ignore regions removed by box center; overlap 0.5.

<sup>c</sup>YOLOX ablation checkpoint released with ByteTrack [3] (trained partly on the first half of each MOT17 training sequence), input 1440×800, confidence 0.01; tracker unchanged. Every tenth frame, the Canny edge fraction of a quarter-scale image divided by 0.14 (capped at 1), averaged over the last seven analyses, sets NMS 0.80 above the selected threshold, else 0.70. HOTA with standard MOT17 pedestrian preprocessing on the validation-half ground truth. Against the default static NMS 0.70 (69.17% HOTA, highest of the static settings 0.70, 0.75, 0.80):  $\Delta$ HOTA −0.07 [−0.54, +0.19], W/T/L 4/0/3.

compared, because neither pipeline has a matched timing on the evaluated split.

## XIV. Limitations

The study has the following limitations.

- 1) *One primary detector and tracker.* All main results use a COCO-pretrained YOLOv8n and ByteTrack. The absolute accuracy is limited by the absence of aerial fine-tuning, and the conclusions about which parameters matter may differ for detectors with different confidence calibration or for trackers with appearance models. Because every confidence threshold used is at or above the tracker’s high threshold, ByteTrack’s low-score association stage received no input in any configuration (Sec. VII); the results characterize its first stage with the tuned buffer and matching threshold, not the full two-stage design. The supplementary pipelines of Sec. XIII reduce this limitation but do not resolve it: their policies differ from Q and from each other, only the U2MOT

comparison is held-out, and its comparator is the author-calibrated static operating point rather than a matched static anchor.

- 2) *Research evaluation protocol.* The VisDrone evaluation is class-agnostic and uses a custom ground-truth filter; it is not the official VisDrone protocol, and the numbers are not comparable with published leaderboard results. The U2MOT experiment uses the protocol of the U2MOT repository, which is not official either.
- 3) *Scope of comparison.* The controller acts on the operating point of a given detector and tracker, so it is compared with static and adaptive settings of the same pipeline rather than with trackers designed specifically for aerial video [13]; such trackers could in principle host the same controller, which was not tested.
- 4) *Processor and timing.* Throughput was measured on a Tesla T4, a data-center graphics processor, not on an embedded flight computer, and excludes JPEG decoding. With decoding, throughput fell below the gate. No flight test was performed. Between-session timing variability is several frames per second.
- 5) *Post hoc attribution.* The matched static anchor was evaluated after the freeze and timed in another session. It is the correct control for the attribution question, but it was not a preregistered competitor. No allocation of the higher resolution at the same average pixel budget as a scene-driven controller was run within this pipeline; the cue audit (Sec. XII) did not match pixel budgets either.
- 6) *Sample size.* All design decisions used seven validation sequences; the weakly identified blur and low-light weights and the transfer failure of the identity-switch constraint are consistent with that size. Intervals come from 5,000 resamples of 17 or 20 sequences (10,000 resamples of 17 or 7 sequences in Sec. XIII); differences of a few tenths of a HOTA point are below the resolution of the test.
- 7) *Provenance.* The notebook cell behind the hand-designed component ablation was not recovered (its outputs and evaluation chain are preserved); the quality-profile test worker resumed from sequence checkpoints; the balanced profile’s first test attempt was interrupted before any result was saved. Exploratory robustness experiments (noise, compression, frame dropping) have no canonical record and are not reported.

## XV. Conclusions

This paper separated the value of calibrating the detector operating point from the value of switching it with the scene, for multi-object tracking in unmanned aerial vehicle video. A scene complexity index built from five calibrated cues drove the confidence threshold, the suppression threshold, and the input resolution of a fixed detector feeding ByteTrack, and all parameters were selected on validation data and frozen before a locked held-out test.

The validation-selected quality profile improved held-out HOTA over a static default by 5.41 points on VisDrone test-dev and by 4.30 points on UAVDT without retuning, and it improved over a hand-designed adaptive controller by 1.14 points while producing significantly more identity switches. A matched static configuration at the profile’s typical operating point reached the same held-out accuracy, so the improvement is due to the validation-calibrated operating

point and tracker profile rather than to frame-to-frame switching. In the primary held-out study and in the separate held-out U2MOT experiment, scene switching did not measurably outperform a well-calibrated static operating point. A separate SparseTrack validation experiment showed a small in-sample HOTA gain over its predeclared static comparator but not over the tracker’s default static setting, suggesting that switching may help relative to some static settings; this does not establish held-out generalization. The balanced profile provided a reproducible trade-off with fewer identity switches, fewer false positives, and higher throughput on VisDrone.

The ablations bound the value of scene-driven switching: it added accuracy for the hand-designed controller, whose resolution range lies on the steep part of the accuracy–resolution curve, although a separate cue audit suggests that this came mainly from its larger average pixel budget, and it added none for the calibrated controller, whose search placed it at the top of that curve within the available throughput budget. For aerial tracking systems, the practical conclusion is that the detector operating point and the tracker profile should be calibrated on representative aerial validation data before any adaptive mechanism is added, and that adaptation should be evaluated against a matched static operating point.

## Acknowledgments

The authors thank the maintainers of the VisDrone and UAVDT benchmarks and of the TrackEval library.

## References

- [1] Bewley, A., Ge, Z., Ott, L., Ramos, F., and Upcroft, B., “Simple online and realtime tracking,” *2016 IEEE International Conference on Image Processing (ICIP)*, IEEE, 2016, pp. 3464–3468. doi:10.1109/icip.2016.7533003.
- [2] Wojke, N., Bewley, A., and Paulus, D., “Simple online and realtime tracking with a deep association metric,” *2017 IEEE International Conference on Image Processing (ICIP)*, IEEE, 2017, pp. 3645–3649. doi:10.1109/icip.2017.8296962.
- [3] Zhang, Y., Sun, P., Jiang, Y., Yu, D., Weng, F., Yuan, Z., Luo, P., Liu, W., and Wang, X., “ByteTrack: Multi-object Tracking by Associating Every Detection Box,” *Computer Vision – ECCV 2022*, Springer Nature Switzerland, 2022, pp. 1–21. doi:10.1007/978-3-031-20047-2\_1.
- [4] Zhu, P., Wen, L., Du, D., Bian, X., Fan, H., Hu, Q., and Ling, H., “Detection and Tracking Meet Drones Challenge,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 44, No. 11, 2022, pp. 7380–7399. doi:10.1109/tpami.2021.3119563.
- [5] Du, D., Qi, Y., Yu, H., Yang, Y., Duan, K., Li, G., Zhang, W., Huang, Q., and Tian, Q., “The Unmanned Aerial Vehicle Benchmark: Object Detection and Tracking,” *Computer Vision – ECCV 2018*, Springer International Publishing, 2018, pp. 375–391. doi:10.1007/978-3-030-01249-6\_23.

- [6] Wu, X., Li, W., Hong, D., Tao, R., and Du, Q., “Deep Learning for Unmanned Aerial Vehicle-Based Object Detection and Tracking: A survey,” *IEEE Geoscience and Remote Sensing Magazine*, Vol. 10, No. 1, 2022, pp. 91–124. doi:10.1109/mgrs.2021.3115137.
- [7] Miller, N., and Rogers, J., “Simultaneous Tracking of Multiple Ground Targets from a Multirotor Unmanned Aerial Vehicle,” *Journal of Aerospace Information Systems*, Vol. 12, No. 3, 2015, pp. 345–364. doi:10.2514/1.i010345.
- [8] Desai, A., and Lee, D.-J., “Efficient Feature Descriptor for Unmanned Aerial Vehicle Ground Moving Object Tracking,” *Journal of Aerospace Information Systems*, Vol. 14, No. 6, 2017, pp. 345–350. doi:10.2514/1.i010503.
- [9] Choi, A. J., Park, J., and Han, J.-H., “Automated Aerial Docking System Using Onboard Vision-Based Deep Learning,” *Journal of Aerospace Information Systems*, Vol. 19, No. 6, 2022, pp. 421–436. doi:10.2514/1.i011053.
- [10] Lu, Y.-R., and Sun, D., “Development of Real-Time Unmanned Aerial Vehicle Urban Object Detection System with Federated Learning,” *Journal of Aerospace Information Systems*, Vol. 21, No. 7, 2024, pp. 547–553. doi:10.2514/1.i011378.
- [11] Parsons, C., Paulson, Z., Nykl, S., Dallman, W., Woolley, B. G., and Pecarina, J., “Analysis of Simulated Imagery for Real-Time Vision-Based Automated Aerial Refueling,” *Journal of Aerospace Information Systems*, Vol. 16, No. 3, 2019, pp. 77–93. doi:10.2514/1.i010658.
- [12] Kanellakis, C., and Nikolakopoulos, G., “Survey on Computer Vision for UAVs: Current Developments and Trends,” *Journal of Intelligent & Robotic Systems*, Vol. 87, No. 1, 2017, pp. 141–168. doi:10.1007/s10846-017-0483-z.
- [13] Liu, S., Li, X., Lu, H., and He, Y., “Multi-Object Tracking Meets Moving UAV,” *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2022, pp. 8866–8875. doi:10.1109/cvpr52688.2022.00867.
- [14] Cheng, G., Yuan, X., Yao, X., Yan, K., Zeng, Q., Xie, X., and Han, J., “Towards Large-Scale Small Object Detection: Survey and Benchmarks,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2023, pp. 1–20. doi:10.1109/tpami.2023.3290594.
- [15] Lin, T.-Y., Dollar, P., Girshick, R., He, K., Hariharan, B., and Belongie, S., “Feature Pyramid Networks for Object Detection,” *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2017, pp. 936–944. doi:10.1109/cvpr.2017.106.
- [16] Kalman, R. E., “A New Approach to Linear Filtering and Prediction Problems,” *Journal of Basic Engineering*, Vol. 82, No. 1, 1960, pp. 35–45. doi:10.1115/1.3662552.
- [17] Kuhn, H. W., “The Hungarian method for the assignment problem,” *Naval Research Logistics Quarterly*, Vol. 2, No. 1-2, 1955, pp. 83–97. doi:10.1002/nav.3800020109.
- [18] Bochinski, E., Eiselein, V., and Sikora, T., “High-Speed tracking-by-detection without using image information,” *2017 14th IEEE International Conference on Advanced Video and Signal Based Surveillance (AVSS)*, IEEE, 2017, pp. 1–6. doi:10.1109/avss.2017.8078516.

- [19] Cao, J., Pang, J., Weng, X., Khirodkar, R., and Kitani, K., “Observation-Centric SORT: Rethinking SORT for Robust Multi-Object Tracking,” *2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2023, pp. 9686–9696. doi:10.1109/cvpr52729.2023.00934.
- [20] Maggolino, G., Ahmad, A., Cao, J., and Kitani, K., “Deep OC-Sort: Multi-Pedestrian Tracking by Adaptive Re-Identification,” *2023 IEEE International Conference on Image Processing (ICIP)*, IEEE, 2023, pp. 3025–3029. doi:10.1109/icip49359.2023.10222576.
- [21] Aharon, N., Orfaig, R., and Bobrovsky, B.-Z., “BoT-SORT: Robust Associations Multi-Pedestrian Tracking,” arXiv preprint arXiv:2206.14651, 2022.
- [22] Du, Y., Zhao, Z., Song, Y., Zhao, Y., Su, F., Gong, T., and Meng, H., “StrongSORT: Make DeepSORT Great Again,” *IEEE Transactions on Multimedia*, Vol. 25, 2023, pp. 8725–8737. doi:10.1109/tmm.2023.3240881.
- [23] Zhang, Y., Wang, C., Wang, X., Zeng, W., and Liu, W., “FairMOT: On the Fairness of Detection and Re-identification in Multiple Object Tracking,” *International Journal of Computer Vision*, Vol. 129, No. 11, 2021, pp. 3069–3087. doi:10.1007/s11263-021-01513-4.
- [24] Zeng, F., Dong, B., Zhang, Y., Wang, T., Zhang, X., and Wei, Y., “MOTR: End-to-End Multiple-Object Tracking with Transformer,” *Computer Vision – ECCV 2022*, Springer Nature Switzerland, 2022, pp. 659–675. doi:10.1007/978-3-031-19812-0\_38.
- [25] Redmon, J., Divvala, S., Girshick, R., and Farhadi, A., “You Only Look Once: Unified, Real-Time Object Detection,” *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2016, pp. 779–788. doi:10.1109/cvpr.2016.91.
- [26] Ren, S., He, K., Girshick, R., and Sun, J., “Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 39, No. 6, 2017, pp. 1137–1149. doi:10.1109/tpami.2016.2577031.
- [27] Bodla, N., Singh, B., Chellappa, R., and Davis, L. S., “Soft-NMS — Improving Object Detection with One Line of Code,” *2017 IEEE International Conference on Computer Vision (ICCV)*, IEEE, 2017, pp. 5562–5570. doi:10.1109/iccv.2017.593.
- [28] Kupperts, F., Kronenberger, J., Shantia, A., and Haselhoff, A., “Multivariate Confidence Calibration for Object Detection,” *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, IEEE, 2020, pp. 1322–1330. doi:10.1109/cvprw50498.2020.00171.
- [29] Akiba, T., Sano, S., Yanase, T., Ohta, T., and Koyama, M., “Optuna: A Next-generation Hyperparameter Optimization Framework,” *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ACM, 2019, pp. 2623–2631. doi:10.1145/3292500.3330701.
- [30] Han, Y., Huang, G., Song, S., Yang, L., Wang, H., and Wang, Y., “Dynamic Neural Networks: A Survey,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Vol. 44, No. 11, 2022, pp. 7436–7456. doi:10.1109/tpami.2021.3117837.

- [31] Jiang, J., Ananthanarayanan, G., Bodik, P., Sen, S., and Stoica, I., “Chameleon: Scalable Adaptation of Video Analytics,” *Proceedings of the 2018 Conference of the ACM Special Interest Group on Data Communication*, ACM, 2018, pp. 253–266. doi:10.1145/3230543.3230574.
- [32] Bernardin, K., and Stiefelhausen, R., “Evaluating Multiple Object Tracking Performance: The CLEAR MOT Metrics,” *EURASIP Journal on Image and Video Processing*, Vol. 2008, 2008, pp. 1–10. doi:10.1155/2008/246309.
- [33] Ristani, E., Solera, F., Zou, R., Cucchiara, R., and Tomasi, C., “Performance Measures and a Data Set for Multi-target, Multi-camera Tracking,” *Computer Vision – ECCV 2016 Workshops*, Springer International Publishing, 2016, pp. 17–35. doi:10.1007/978-3-319-48881-3\_2.
- [34] Luiten, J., Ošep, A., Dendorfer, P., Torr, P., Geiger, A., Leal-Taixé, L., and Leibe, B., “HOTA: A Higher Order Metric for Evaluating Multi-object Tracking,” *International Journal of Computer Vision*, Vol. 129, No. 2, 2021, pp. 548–578. doi:10.1007/s11263-020-01375-2.
- [35] Dendorfer, P., Ošep, A., Milan, A., Schindler, K., Cremers, D., Reid, I., Roth, S., and Leal-Taixé, L., “MOTChallenge: A Benchmark for Single-Camera Multiple Target Tracking,” *International Journal of Computer Vision*, Vol. 129, No. 4, 2021, pp. 845–881. doi:10.1007/s11263-020-01393-0.
- [36] Efron, B., and Tibshirani, R. J., *An Introduction to the Bootstrap*, Springer US, 1993. doi:10.1007/978-1-4899-4541-9.
- [37] Pech-Pacheco, J., Cristobal, G., Chamorro-Martinez, J., and Fernandez-Valdivia, J., “Diatom autofocusing in brightfield microscopy: a comparative study,” *Proceedings 15th International Conference on Pattern Recognition. ICPR-2000*, IEEE Comput. Soc, 2000, pp. 314–317. doi:10.1109/icpr.2000.903548.
- [38] Zhao, Y., Lv, W., Xu, S., Wei, J., Wang, G., Dang, Q., Liu, Y., and Chen, J., “DETRs Beat YOLOs on Real-time Object Detection,” *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2024, pp. 16965–16974. doi:10.1109/cvpr52733.2024.01605.
- [39] Liu, K., Jin, S., Fu, Z., Chen, Z., Jiang, R., and Ye, J., “Uncertainty-aware Unsupervised Multi-Object Tracking,” *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2023, pp. 9962–9971. doi:10.1109/iccv51070.2023.00917.
- [40] Ge, Z., Liu, S., Wang, F., Li, Z., and Sun, J., “YOLOX: Exceeding YOLO Series in 2021,” arXiv preprint arXiv:2107.08430, 2021.
- [41] Liu, Z., Wang, X., Wang, C., Liu, W., and Bai, X., “SparseTrack: Multi-Object Tracking by Performing Scene Decomposition Based on Pseudo-Depth,” *IEEE Transactions on Circuits and Systems for Video Technology*, Vol. 35, No. 5, 2025, pp. 4870–4882. doi:10.1109/tcsvt.2024.3524670.
- [42] Milan, A., Leal-Taixe, L., Reid, I., Roth, S., and Schindler, K., “MOT16: A Benchmark for Multi-Object Tracking,” arXiv preprint arXiv:1603.00831, 2016.