

Training-Free Self-Calibrating Control Layer for Heterogeneous Multi-Object Trackers

Ahmed Gouda Ismail^{1*}, Mohamed S. Mohamed¹ and Tarek Ahmed Mahmoud²

^{1*}Computer Engineering and Artificial Intelligence Department, Military Technical College, Cairo, Egypt.

²Computer Engineering Department, Egypt University of Informatics, Cairo, Egypt.

*Corresponding author(s). E-mail(s): a7medgouda1@gmail.com;
Contributing authors: mohamedms@mtc.edu.eg; tarek.mohamed@eui.edu.eg;

Abstract

Multi-object tracking pipelines are often built from a detector and a tracker developed separately. When the detector is replaced, retrained, or quantized for embedded hardware, its confidence scale changes, while the tracker keeps thresholds tuned for another detector. This paper presents SCL, a training-free self-calibrating control layer between a frozen detector and a frozen tracker. It calibrates itself online with nested Otsu thresholds on the detector logits of recent frames and decides whether the score stream is clean or noisy. On clean streams it keeps the tracker's operating point. On noisy streams it selects the candidates that reach the tracker, remaps their scores, and sets the tracker's thresholds. It also gives a continuation stage to trackers that lack one. It needs no labels and uses no future frames. One frozen configuration was evaluated on seven published trackers, four detection sources, and two datasets with paired bootstrap intervals. It raised the higher order tracking accuracy (HOTA) of OC-SORT and of the external PD-SORT by 0.61 points, with intervals above zero. When the detector scores were halved, it restored two trackers from 0 to about 66 HOTA. With a noisy RT-DETR-L detector, it raised the multi-object tracking accuracy (MOTA) of ByteTrack and BoT-SORT by 6.0 and 7.9 points. On four trackers with a low-score stage and well-matched detections, HOTA changed by at most 0.06 points. SCL takes 2.95 ms per frame on a four-core CPU without a GPU, so it adds very little latency and is suitable for real-time use.

Keywords: Multi-object tracking, Tracking by detection, Detector confidence calibration, Online adaptation, Otsu thresholding, Real-time image processing

1 Introduction

Multi-object tracking (MOT) finds every object of interest in a video and keeps the same identity for each object over time. It is a core part of surveillance, driving, and airborne systems [1–3], and many of these systems must produce results with low delay. Most practical trackers follow the tracking-by-detection design. A detector scores

candidate boxes in each frame, and the tracker links the boxes whose scores pass its thresholds into trajectories [4–6].

In deployed systems the detector and the tracker are rarely developed together. A detector is replaced by a faster model, retrained for a new sensor, or quantized for embedded hardware. The tracker, however, keeps the thresholds

that its authors tuned for another detector. Detector confidence is not calibrated across models and conditions [7], so the new detector usually has a different score scale, and nothing in the pipeline notices the change.

The effect can be large. In the experiments of this paper, halving the scores of a published detector drives two well-known trackers to 0 higher order tracking accuracy (HOTA) [8], because no detection reaches their fixed thresholds (Sect. 6.3). Retuning every tracker for every detector needs labelled data from the deployment domain, which an operational system usually does not have.

Existing methods only partly address this problem. Adaptive-threshold trackers adjust their own thresholds from the detections of each frame, but each method is built into one specific tracker [9–11]. Confidence calibration methods correct the detector scores, but they need labelled data [7]. Adaptive inference methods learn or profile a policy for one model [12, 13]. What is missing is a training-free component that works outside the tracker, calibrates itself from the detector’s own score stream, and leaves the tracker unchanged when the tracker already fits that stream.

This paper presents such a component: SCL, a self-calibrating control layer between a frozen detector and a frozen tracker. It reads the stream of detector scores and the operating point that the tracker declares, and it decides online whether the tracker needs help. SCL is described in Sect. 3, and Sect. 4 explains how it was developed from an earlier design that always intervened.

The main contributions are:

1. SCL, a training-free control layer that calibrates itself from the detector score stream and changes the tracker’s input only when the stream is noisy (Sect. 3).
2. A host contract of four numbers that lets one frozen configuration work with single-stage and two-stage trackers without modifying them (Sect. 3.6).
3. A development path from always-on banding (AOB), a prior design that always intervened, in which each step is compared with that baseline (Sect. 4).
4. An evaluation of one frozen configuration on seven published trackers, four detection sources, and two datasets, with development

and predeclared external roles, paired bootstrap intervals, and a latency measurement on a CPU (Sects. 5 and 6).

The rest of the paper is organized as follows. Section 2 reviews related work. Section 3 describes SCL, and Section 4 explains how it was developed. Section 5 gives the experimental setup, Section 6 the results, and Section 7 a discussion. Section 8 concludes the paper.

2 Related Work

2.1 Tracking by detection

SORT predicts each box with a Kalman filter and matches the predictions to new detections by overlap with the Hungarian method [4]. StrongSORT and BoT-SORT add appearance features [14, 15]. ByteTrack introduced a second association stage that uses low-confidence detections to continue existing tracks [5]. This two-stage design was kept by BoT-SORT [15], SparseTrack [16], BoostTrack [17], and Hybrid-SORT [18]. OC-SORT [6], Deep OC-SORT [19], and PD-SORT [20] use a single confidence threshold. Recent designs learn association from box coordinates (C-TWiX [21]), run parallel association paths (TOPICTrack [22]), or associate two detection views (TrackTrack [23]). End-to-end trackers predict identities directly [24, 25]. All of these trackers read detector scores through fixed thresholds that were tuned for one detector, usually YOLOX [26]. Seven of them are used in this paper as hosts and are kept unchanged.

2.2 Detector confidence and calibration

Detectors of different families, such as YOLOX [26], YOLOv8 [27], YOLOv9 [28], DETR [29], and RT-DETR [30], produce different score distributions. Their confidence is not calibrated across models and conditions, and calibration methods for object detection fit a mapping on labelled data [7]. Otsu’s method selects a threshold from a histogram without labels by maximizing the separation between two classes [31]. SCL applies it to the detector logits of recent frames.

Table 1 Adaptive methods closest to SCL (BT: ByteTrack).

Method	Adapted	Statistics	Where
Adapt. BT [9]	confidence threshold	current frame	inside BT
AdapTrack [10]	matching thresholds	current frame	inside the tracker
ACT [11]	high and low thresholds	current frame	inside the tracker
Küppers [7]	detector scores	labelled data	detector output
Chameleon [13]	pipeline configuration	profiling	analytics pipeline
SCL (this paper)	candidates, scores, thresholds	past frames, no labels	between detector and tracker

2.3 Adaptive thresholds inside trackers

Several trackers adapt their own confidence thresholds. Adaptive ByteTrack derives the confidence threshold of ByteTrack from the detections of each frame [9]. AdapTrack adapts the thresholds of its matching step [10]. ACT sets the high and low thresholds of each frame from the first quartile of the confidence scores and the number of candidates [11]. These methods are built into one tracker and use statistics of the current frame. SCL differs in four ways. Its statistics come from past frames. It decides for the whole stream whether to intervene at all, and otherwise keeps the tracker’s own operating point. It acts through a declared host contract on several published trackers without modifying them. It is evaluated as one frozen policy under confidence shift and on external systems.

2.4 Adaptive inference

Dynamic neural networks adapt their computation to each input [12], and Chameleon re-profiles the configuration of video analytics pipelines over time [13]. These methods learn or profile a policy for each model. A controller that sets the detector operating point from image cues is also tied to the detector it was calibrated with, because a change of the detector’s score scale leaves the image unchanged. SCL reads the score stream instead and acts on the tracker’s input and thresholds. Table 1 summarizes the differences.

3 Method

3.1 Problem setting

At frame t a frozen detector emits candidates $D_t = \{(\mathbf{x}_i, s_i, c_i)\}$ with boxes $\mathbf{x}_i$, scores $s_i \in (0, 1)$, and class labels c_i , down to an emission floor. The layer works on the logit of each score,

$$z_i = \log \frac{s_i}{1 - s_i}. \quad (1)$$

A frozen tracker, called the host, declares a host contract $h = (a, b, \ell, m_0)$. Here a is the score threshold of its first association, b is the lowest score from which a new track can start, ℓ is the lowest score that any of its association stages uses, and m_0 is its matching tolerance (one minus its minimum overlap). A single-stage tracker has $\ell = \min(a, b)$, and a two-stage tracker has $\ell < \min(a, b)$.

For every frame the layer outputs the set K_t of candidates passed to the tracker, the scores $\tilde{s}$ passed with them, and the thresholds (a_t, b_t) and tolerance m_t that the tracker uses in that frame:

$$(K_t, \tilde{s}, a_t, b_t, m_t) = \mathcal{L}(Z_{<t}, D_t, \mathcal{T}_{t-1}, g_t; h). \quad (2)$$

$Z_{<t}$ are the logits of earlier frames, $\mathcal{T}_{t-1}$ are the tracks that the tracker reported for frame $t - 1$, and g_t is an optional image-motion cue of frame t . The layer works online. Its statistics use only frames before t , and the data of frame t enter its state after the decision. No future frame is used.

The layer has four requirements: it needs no training and no labelled calibration data; one frozen configuration serves every tracker, detector, and dataset; the output stays unchanged when the tracker’s operating point already fits the score stream; and its cost per frame is small next to the detector.

3.2 Overview

Figure 1 shows SCL. It has four parts: stream statistics, a regime estimate, a candidate policy for each regime, and a continuation rescue. The host contract enters the rescue and the thresholds passed to the tracker. The detector and the tracker are not modified.




Fig. 1 Self-calibrating control layer (SCL) between a frozen detector and a frozen tracker. The statistics use earlier frames only. The host contract is the tracker’s published operating point. The tracker receives the selected candidates, their scores, and its thresholds for the frame.

3.3 Self-calibration from the score stream

The layer pools the logits of the candidates of the previous $W = 10$ frames and splits them twice with the exact Otsu threshold [31]. The first split t_1 separates background from foreground. A second split of the foreground, t_2 , separates ambiguous from confident candidates. The confident share of the foreground,

$$\rho_t = \frac{|\{i : z_i \geq t_2\}|}{|\{i : z_i \geq t_1\}|}, \quad (3)$$

summarizes how cleanly the detector separates objects from clutter. Because the splits are computed on logits, they follow any increasing affine change of the logit, such as a change of temperature, and the bands keep the same candidates.

The bands have their intended meaning only if the stream has a background mode. When the class below t_1 holds fewer pooled candidates than the foreground, which happens when a detector emits few low-scoring candidates, the first split falls inside the objects. In that case the frame counts as clean evidence, $\rho_t = 1$. This background check was added after a development case in which sequence starts with about ten candidates per frame produced splits at $t_1 \approx 0.6$ and $t_2 \approx 0.93$ in probability terms and a wrong noisy decision.

3.4 Intervention regimes

Regime estimate. The stream is clean at frame t if the median $\bar{\rho}_t$ of ρ over all non-cold frames so far is at least one half, and noisy otherwise. Before the first valid split, for example at the first frame of a stream, the regime is cold. The regime is therefore

a property of the detector and scene stream, and a short dip in confidence, for example under fast camera motion, does not flip it.

Cold and clean regimes. The tracker’s own thresholds are kept. In the clean regime they are lowered only as far as t_2 when the tracker would otherwise reject the confident class:

$$a_t = \min(a, \sigma(t_2)), \quad b_t = \min(b, \sigma(t_2)), \quad (4)$$

where σ is the logistic function. Scores are passed unchanged, and only cross-class duplicates, that is, one box proposed with two class labels, are removed. On a stream whose confident class lies above the tracker’s thresholds, the tracker therefore receives exactly its native input.

Noisy regime. The candidates are split into bands. The primary band $z \geq t_2$ may start and continue tracks, the extension band $t_1 \leq z < t_2$ may only continue tracks, and candidates below t_1 are discarded. Scores are remapped through their rank u_i in the empirical distribution of recent stream scores,

$$\tilde{s}_i = \begin{cases} 0.5 + 0.5 u_i, & z_i \geq t_2, \\ 0.1 + 0.4 u_i, & t_1 \leq z_i < t_2, \end{cases} \quad (5)$$

and the tracker’s association and birth thresholds are set to $a_t = b_t = 0.5$. Extension candidates can therefore reach only a tracker’s low-score stage. Duplicates are resolved with track context: a weaker candidate that overlaps a stronger one with an intersection over union (IoU) above one half is removed, unless a different track of frame $t - 1$ claims it. This keeps occluded people in crowds.

3.5 Continuation rescue and motion rule

A foreground candidate ($z \geq t_1$, or any emitted candidate for a single-stage tracker when the stream has no background mode) may continue an existing track of frame $t-1$ with IoU of at least one half while no usable candidate covers that track and while the tracker cannot see the candidate at the operating point passed for this frame. Such a candidate is passed with its score raised just above the lowest threshold that the tracker uses in that frame. For a two-stage tracker the rescue can act only on foreground candidates below the tracker’s low threshold ℓ . For a single-stage tracker it acts on foreground candidates below the lowest passed threshold, so it supplies the continuation stage that the tracker lacks.

In noisy frames only, and only when an image-motion cue is available, the matching tolerance widens with the ratio r_t of the current global image motion to its past median. The motion cue is the phase-correlation shift [32] between consecutive grayscale frames at quarter resolution, divided by the image diagonal. The tolerance is

$$m_t = \min\left(0.95, 1 - \frac{1 - m_0}{\max(1, r_t)}\right). \quad (6)$$

3.6 Host contract and adapters

Table 2 lists the contracts used in this paper. Each value is a documented property of the published tracker configuration, not a fitted quantity. An adapter for each tracker family maps the decision of the layer onto the tracker’s interface. The selected candidates and their scores replace the detector output of the frame, and the thresholds and tolerance replace the tracker’s own for that frame. Post-processing that the authors apply after tracking, such as interpolation or offline linking, is kept unchanged. The policy and all of its constants are identical for every tracker. A tracker enters only through its published contract values, and the adapter translates the decision without making one.

3.7 Frozen constants

Table 3 lists every constant of the layer. The same values are used for every tracker, detector, and dataset. The band thresholds t_1 and t_2 are not

constants: they are recomputed online for every frame. The duplicate and rescue overlap of one half follows the matching rule of MOT evaluation, in which two boxes with IoU above 0.5 cannot both match distinct objects. The remap ranges of Eq. (5) place the two bands on the common association and birth threshold of 0.5 and the low-score threshold of 0.1. The cap of 0.95 keeps a minimum overlap of 0.05 for every match. The window, the motion history, and the remap memory are declared defaults that were kept fixed in all development stages.

3.8 Per-frame procedure and cost

Algorithm 1 gives the procedure for one frame. Per frame the layer sorts the pooled logits of at most W frames for the two Otsu splits, computes candidate overlaps with the other candidates and with the tracks of the previous frame, and, when the motion cue is used, computes one phase correlation on a quarter-resolution image. With n candidates per frame and k tracks, the cost is $O(Wn \log(Wn) + n^2 + nk)$ plus the fixed image cost. The state is a ring buffer of W frames of logits, the per-frame ρ values, the remap memory, and a history of 100 motion values.

4 Development from a Prior Design

4.1 Baseline: always-on banding

SCL was developed from an earlier design of the same research program, called always-on banding (AOB). AOB computed the same nested Otsu bands but treated every stream as noisy. In every frame it applied the bands, removed every weaker candidate that overlapped a stronger one with IoU above one half, and remapped all scores by rank onto fixed ranges. It admitted no candidate at the first frame. AOB was designed for noisy detector streams, and it is the baseline of this section.

Table 4 compares AOB and SCL with the same strong published trackers on MOT17 [33]. AOB lowered the HOTA of SparseTrack by 4.15 points and that of BoostTrack by 5.89 points online and 5.37 points with interpolation, all with intervals below zero. Two causes were found. First, overlap-based duplicate removal deleted occluded people in crowds. Second, the self-calibrated thresholds

Table 2 Host contracts: the operating point that each published tracker declares, as used by the layer. τ is the tracker’s own published threshold, which differs between trackers.

Tracker (setting)	Type	assoc a	birth b	low ℓ	match m_0	Decision mapped to
ByteTrack (official MOT17)	two-stage	0.6	0.7	0.1	0.8	thresholds, candidates, scores
ByteTrack (library default)	two-stage	0.25	0.25	0.1	0.8	same
SparseTrack	two-stage	τ	$\tau + 0.1$	0.1	published	same
BoostTrack	two-stage	τ	τ	0.1	1-IoU gate	same
Hybrid-SORT	two-stage	τ	τ	0.1	1-IoU gate	same
OC-SORT	single-stage	0.6	0.6	0.6	0.7	same
PD-SORT	single-stage	τ	τ	τ	1-IoU gate	same

Table 3 Frozen constants of SCL.

Constant	Value
Statistics window W	10 frames
Regime rule	clean if median $\rho \geq 0.5$ over non-cold frames so far
Background check	$\rho = 1$ if fewer candidates below t_1 than above
Noisy score ranges	primary $[0.5, 1)$, extension $[0.1, 0.5)$
Noisy association and birth thresholds	0.5
Remap memory	20 frames, one every 10 frames
Duplicate and rescue IoU	0.5
Rescued score	lowest used threshold + 0.001
Motion cue	phase correlation, quarter resolution
Motion history, warm-up	100, 5 frames
Tolerance cap	0.95

were stricter than the trackers’ own: for SparseTrack and online BoostTrack, false positives fell by 1,630 and 1,351, but false negatives rose by 4,955 and 6,167. On the same trackers, SCL changed HOTA by +0.06 for SparseTrack and left both BoostTrack variants identical to the host. Figure 2 shows the same comparison.

4.2 Design steps

Each step from AOB to SCL was introduced to remove one diagnosed cause and was compared with the step before it on development data (Table 5). Emulating AOB inside the new framework reduced ByteTrack at floor 0.1 from 67.70 to 56.41 HOTA and OC-SORT from 66.44 to 45.72. The regime gate, which adds the regime estimate, the self-limiting clean regime, and the crowd-safe duplicate rule, restored SparseTrack from 64.72 to 68.93 and BoostTrack from 62.61 to 68.37. Estimating the regime from the whole stream removed regime changes in the middle of a stream and

Algorithm 1 SCL for frame t

Require: candidates D_t , tracks $\mathcal{T}_{t-1}$, cue g_t , contract (a, b, ℓ, m_0)

- 1: $z_i \leftarrow$ logit of s_i for every candidate
- 2: **if** no valid split exists **then**
- 3: regime $\leftarrow$ cold; $(a_t, b_t) \leftarrow (a, b)$
- 4: **else**
- 5: $(t_1, t_2) \leftarrow$ nested Otsu on the window
- 6: $\rho_t \leftarrow$ Eq. (3) with the background check
- 7: regime $\leftarrow$ clean if median(ρ) ≥ 0.5 , else noisy
- 8: **end if**
- 9: **if** regime is cold or clean **then**
- 10: remove cross-class duplicates; keep raw scores
- 11: **if** clean **then** set (a_t, b_t) by Eq. (4)
- 12: **end if**
- 13: **else**
- 14: remove duplicates with track context
- 15: drop $z < t_1$; remap scores by Eq. (5)
- 16: $a_t \leftarrow 0.5$; $b_t \leftarrow 0.5$
- 17: **end if**
- 18: rescue track continuations (not in cold frames)
- 19: $m_t \leftarrow$ Eq. (6) if noisy, else m_0
- 20: pass $K_t, \tilde{s}, a_t, b_t, m_t$ to the tracker
- 21: add the logits of frame t to the window; add g_t to the history
- 22: store the tracks of frame t after the tracker update

raised ByteTrack from 67.561 to 67.702 and OC-SORT from 65.697 to 66.071. The continuation rescue added 0.2 to 0.4 points for OC-SORT. The background check completed SCL: it returned ByteTrack and BoostTrack to output identical to the host and raised OC-SORT to 67.04 and 66.91 at the two floors. Four variants were rejected on evidence, because they lowered HOTA, added

Table 4 Baseline comparison of the development stage on MOT17 val-half (published YOLOX-X detections). Host: published tracker alone (our reproduction). Δ : tracker with the prior design AOB or with SCL minus host, with 95% paired-bootstrap interval (10,000 resamples, seed 42). Bold intervals exclude zero.

Metric	Host	Δ AOB	Δ SCL
<i>SparseTrack</i> [16]			
HOTA	68.88	-4.15	+0.06
		[-5.54, -1.66]	[-0.01, +0.25]
MOTA	77.85	-6.14	+0.08
		[-8.10, -2.74]	[-0.03, +0.30]
IDF1	81.97	-4.49	+0.16
		[-6.19, -1.96]	[-0.02, +0.70]
<i>BoostTrack, online</i> [17]			
HOTA	68.49	-5.89	0 (identical)
		[-7.76, -2.87]	
MOTA	75.50	-8.87	0 (identical)
		[-12.37, -7.14]	
IDF1	81.41	-6.22	0 (identical)
		[-9.33, -2.85]	
<i>BoostTrack, with interpolation</i> [17]			
HOTA	71.72	-5.37	0 (identical)
		[-7.30, -2.75]	
MOTA	81.03	-8.73	0 (identical)
		[-12.59, -6.50]	
IDF1	84.16	-5.65	0 (identical)
		[-8.96, -2.77]	




Fig. 2 HOTA on MOT17 val-half of three strong published trackers alone, with AOB, and with SCL. The labels give the change caused by AOB.

false positives, or let false tracks through on a noisy stream.

4.3 Freeze and tests

The final configuration was frozen before any external evaluation. It was recorded in a version-control tag and protected by a lock file that holds

cryptographic hashes of the layer, its configuration, and the evaluation runners. Every later run verified the lock before running. Sixty-one automated tests check that the layer uses no future data, that its state resets between sequences, that clean streams pass through unchanged, that the output does not depend on detector, tracker, or sequence names, that affine score transforms leave the bands unchanged where designed, and that the lock is intact.

5 Experimental Setup

5.1 Hosts and baselines

Seven published trackers serve as hosts: ByteTrack [5], BoT-SORT [15], OC-SORT [6], SparseTrack [16], BoostTrack [17], PD-SORT [20], and Hybrid-SORT [18]. Each host uses the authors' code and published configuration, and its contract is given in Table 2.

In every comparison the baseline is the host alone: the published tracker run by us with its published configuration and the same detections. The second arm adds the frozen SCL and changes nothing else. Both arms run in the same environment. Where the authors report their own numbers, these are listed next to our reproduction of the host.

5.2 Detectors and data

Table 6 lists the evaluation cells. On MOT17 val-half [33] the trackers use the published YOLOX-X detections [26]. The emission floor, the lowest detector score passed on, is 0.01 or 0.1 (Table 7). On KITTI tracking training [34], the trackers run with two detectors that differ from the MOT17 detector: YOLOv8n [27] and RT-DETR-L [30]. To isolate the calibration problem, the published MOT17 detections were also transformed before the tracker by four monotone maps: cubing (s^3), halving ($0.5s$), and logit temperatures $T = 2$ and $T = 0.5$, $\sigma(z/T)$. The trackers kept their published thresholds. The external systems use the authors' released detections on MOT17 val-half.

5.3 Evaluation roles

Every cell has one of two roles (Table 6). Development cells were visible while the layer was designed: MOT17 with four trackers, KITTI with

Table 5 Design steps from always-on banding (AOB) to SCL on development data (HOTA, %), and variants that were tested and rejected. The floor of ByteTrack and OC-SORT is given in parentheses.

Step	Change	Effect
AOB emulated	always-noisy bands, overlap deduplication, rank remap	ByteTrack (0.1) 56.41 vs host 67.70; OC-SORT (0.1) 45.72 vs 66.44
+ regime gate	regime statistic, host-anchored clean regime, crowd-safe duplicates	SparseTrack 64.72 to 68.93 (host 68.88); BoostTrack 62.61 to 68.37 (host 68.49)
+ whole-stream regime	regime from all past frames	ByteTrack (0.1) 67.561 to 67.702; OC-SORT (0.1) 65.697 to 66.071
+ continuation rescue	foreground track-consistent rescue	OC-SORT (0.01) 66.405 to 66.619; OC-SORT (0.1) 66.071 to 66.469
+ background check = SCL	no background mode gives no evidence against the host	ByteTrack/BoostTrack (0.1) identical to host; OC-SORT 67.041 (0.01), 66.907 (0.1)
<i>Rejected variant</i> no candidates at frame 1 rescue below a two-stage host's low stage rank-remapped scores in clean frames host-clipped noisy primary band		<i>Evidence</i> −0.3 to −0.5 HOTA on every MOT17 host FP +700 to +1000, −0.4 to −0.5 HOTA ByteTrack −0.2 / −1.6 HOTA KITTI YOLOv8n ByteTrack +1.4 HOTA but uncontrolled false tracks on a noisy aerial stream (26.9 vs 14.9 boxes/frame)

Table 6 Evaluation cells and their role. Development cells were visible while the layer was designed; predeclared cells were run once with the frozen policy.

Data	Detector	Trackers	Role
MOT17 val-half (7 sequences)	published YOLOX-X detections	ByteTrack (2 settings), OC-SORT, BoostTrack, SparseTrack	development
KITTI tracking training (21 sequences)	RT-DETR-L; YOLOv8n	ByteTrack and BoT-SORT (RT-DETR-L); OC-SORT (YOLOv8n)	development (detector transfer)
MOT17 val-half, score transforms	published YOLOX-X, transformed scores	ByteTrack (2 settings), OC-SORT, BoostTrack	development (robustness)
MOT17 val-half	released PD-SORT / Hybrid-SORT detections	PD-SORT, Hybrid-SORT	predeclared external

two other detectors, and the score transforms. The predeclared external systems, PD-SORT and Hybrid-SORT, were selected and declared before they were run, and they were run once with the frozen configuration. For each external host, the reproduction against the authors' number is classed as EXACT (within 0.2 HOTA) or CLOSE (within 1.0 HOTA with a documented, untuned cause).

5.4 Metrics and statistics

Results are reported as HOTA [8], MOTA [35], and IDF1 [36]. They are computed with the TrackEval library (commit 12c8791) and the official ground truth of each benchmark, following MOTChallenge practice [37]. On KITTI tracking, HOTA is the official value averaged over cars and pedestrians.

All comparisons are paired. Differences are pooled over sequences, and their 95% intervals come from a percentile bootstrap over sequences [38] with 10,000 resamples and seed 42. A difference is called significant when its interval excludes

zero. As a distribution-free complement, per-sequence wins, ties ($|\Delta\text{HOTA}| < 0.01$), and losses are also counted.

5.5 Latency measurement

Latency was measured after the freeze on one cloud machine with a 4-vCPU Intel Xeon processor at 2.10 GHz, 15 GB of memory, and no GPU. A live YOLOv8n detector (32-bit floats, 736 pixels, batch 1) feeds ByteTrack on three KITTI tracking sequences (0001, 0009, and 0019). Both arms run on the same frames back to back, and the first ten frames of each sequence are excluded, which leaves 2,279 timed frames per arm. The time of the layer includes the motion cue, the per-frame step, and the storage of the tracker output.

6 Results

Figure 3 summarizes the HOTA difference between each host with SCL and the host alone for every evaluated cell. The following subsections




Fig. 3 Difference in HOTA (host with SCL minus host alone) with 95% paired-bootstrap interval for every evaluated cell. Blue: interval above zero; gray: interval contains zero, or identical output.

give the baseline comparison for each part of the evaluation.

6.1 Tracker transfer

Table 7 compares the development trackers with their hosts on MOT17. The two-stage trackers with well-matched detections stayed at the accuracy of their hosts. BoostTrack, online and with interpolation, and ByteTrack at floor 0.1 produced identical output. ByteTrack at floor 0.01 changed by -0.014 HOTA and SparseTrack by $+0.055$. This is the self-limiting behavior of the clean regime (Sect. 3.4), which the prior design did not have (Table 4).

The single-stage OC-SORT improved at both emission floors. Its HOTA rose by 0.61 points ($[+0.36, +1.24]$) at floor 0.01 and by 0.46 points ($[+0.27, +1.09]$) at floor 0.1, and its MOTA by 1.27 and 1.26 points. HOTA was higher on 7 of 7 sequences in both settings (Fig. 4). The gain comes from the continuation rescue. OC-SORT discards every detection below its single threshold, and the layer passes it the foreground candidates that continue a track that would otherwise stay unmatched.

Table 7 Baseline comparison of the development trackers on MOT17 val-half (published YOLOX-X detections). Paper: number reported by the tracker’s authors. Host: published tracker alone (our reproduction). Δ : host with SCL minus host, with 95% paired-bootstrap interval (10,000 resamples, seed 42); bold intervals exclude zero. The floor is the lowest detector score emitted.

Metric	Paper	Host	Δ [95% CI]
<i>SparseTrack</i> [16]			
HOTA	69.2	68.88	+0.06 [−0.01, +0.25]
MOTA	76.8	77.85	+0.08 [−0.03, +0.30]
IDF1	81.4	81.97	+0.16 [−0.02, +0.70]
<i>BoostTrack, online</i> [17]			
HOTA	68.37	68.49	0 (identical)
MOTA	75.56	75.50	0 (identical)
IDF1	81.35	81.41	0 (identical)
<i>BoostTrack, + GBI</i> [17]			
HOTA	71.33	71.72	0 (identical)
MOTA	80.55	81.03	0 (identical)
IDF1	83.84	84.16	0 (identical)
<i>ByteTrack, floor 0.01</i> [5]			
HOTA	–	67.70	−0.014 [−0.058, +0.009]
MOTA	76.6	77.60	+0.06 [−0.15, +0.34]
IDF1	79.3	79.47	−0.031 [−0.135, +0.035]
<i>ByteTrack, floor 0.1</i> [5]			
HOTA	–	67.70	0 (identical)
MOTA	76.6	77.60	0 (identical)
IDF1	79.3	79.47	0 (identical)
<i>OC-SORT, floor 0.01</i> [6]			
HOTA	66.5	66.43	+0.61 [+0.36 , +1.24]
MOTA	74.9	74.67	+1.27 [+0.17 , +3.01]
IDF1	77.7	78.05	+0.66 [−0.25, +1.60]
<i>OC-SORT, floor 0.1</i> [6]			
HOTA	66.5	66.44	+0.46 [+0.27 , +1.09]
MOTA	74.9	74.67	+1.26 [+0.37 , +3.34]
IDF1	77.7	78.05	+0.287 [+0.003 , +0.838]




Fig. 4 Per-sequence HOTA differences (host with SCL minus host alone) on MOT17 val-half for the development and predeclared trackers.

6.2 Detector change

Table 8 reports the KITTI results with the two other detectors. With the noisy RT-DETR-L stream, the two-stage trackers gained in MOTA and IDF1. ByteTrack gained 5.98 MOTA points ($[+1.16, +9.94]$), 3.87 IDF1 points, and

Table 8 Baseline comparison under a detector change on KITTI tracking training (official HOTA, car and pedestrian averaged). Host: published tracker alone with the same detections. Δ : host with SCL minus host, with 95% paired-bootstrap interval; bold intervals exclude zero. BoT-SORT uses 20 of the 21 sequences, because the unmodified tracker gives no valid output on one sequence.

Metric	Host	Δ [95% CI]
<i>ByteTrack</i> [5], <i>RT-DETR-L</i>		
HOTA	50.73	+0.91 [+0.22 , +1.55]
MOTA	47.27	+5.98 [+1.16 , +9.94]
IDF1	64.98	+3.87 [+1.67 , +5.06]
<i>BoT-SORT</i> [15], <i>RT-DETR-L</i>		
HOTA	53.91	+0.61 [−0.39, +1.61]
MOTA	48.25	+7.86 [+2.21 , +13.44]
IDF1	67.74	+3.10 [+1.00 , +4.57]
<i>OC-SORT</i> [6], <i>YOLOv8n</i>		
HOTA	36.27	+7.91 [+5.47 , +9.98]
MOTA	33.99	+9.40 [+1.10 , +12.30]
IDF1	50.46	+9.60 [+5.61 , +12.16]

0.91 HOTA points ([+0.22, +1.55]). BoT-SORT gained 7.86 MOTA points ([+2.21, +13.44]) and 3.10 IDF1 points. With YOLOv8n, the single-stage OC-SORT gained 7.91 HOTA points ([+5.47, +9.98]), 9.40 MOTA points, and 9.60 IDF1 points.

6.3 Detector confidence shift

Table 9 and Fig. 5 give the trackers under the four score transforms. Halving the scores drove ByteTrack (official setting) and OC-SORT to 0 HOTA, because no detection reached their thresholds. With the layer they reached 66.17 and 66.72. Cubing reduced them to 60.92 and 58.34, and the layer restored them to 65.81 and 66.60. BoostTrack rose from 61.33 to 64.84 under cubing. With temperature $T = 2$, the layer raised ByteTrack and OC-SORT by 1.11 and 1.24 points. The library setting of ByteTrack, whose low thresholds already admit the shifted scores, stayed within 0.14 points in all four transforms. Over the 14 conditions, the layer raised HOTA significantly in 7 and lowered it significantly in none.

6.4 External systems

Table 10 gives the two predeclared external systems. PD-SORT [20] was reproduced identically to the authors’ released output. With SCL, its

Table 9 Baseline comparison under detector confidence shift on MOT17 val-half (HOTA). The detector scores are transformed before the tracker, whose published thresholds stay fixed; T is a logit temperature, $\sigma(z/T)$. Host: published tracker alone. Δ : with SCL minus host, with 95% paired-bootstrap interval (10,000 resamples, seed 42). Bold intervals exclude zero. Emission floor 0.01 (BoostTrack 0.1).

Transform	Host	+ SCL	Δ HOTA
<i>ByteTrack</i> , <i>official setting</i> [5]			
s^3	60.92	65.81	+4.89 [+2.60 , +13.69]
0.5 s	0.00	66.17	+66.17 [+49.32 , +75.40]
$T = 2$	65.84	66.95	+1.11 [+0.35 , +3.87]
$T = 0.5$	67.30	67.29	−0.009 [−0.051, +0.005]
<i>ByteTrack</i> , <i>library setting</i> [5]			
s^3	66.31	66.31	−0.006 [−0.384, +0.383]
0.5 s	66.61	66.48	−0.14 [−0.66, +0.14]
$T = 2$	63.97	63.97	0 (identical)
$T = 0.5$	66.53	66.52	−0.006 [−0.029, +0.001]
<i>OC-SORT</i> [6]			
s^3	58.34	66.60	+8.26 [+5.86 , +16.81]
0.5 s	0.00	66.72	+66.72 [+50.32 , +75.56]
$T = 2$	65.99	67.22	+1.24 [+0.52 , +3.42]
$T = 0.5$	66.61	67.06	+0.45 [−0.01, +0.97]
<i>BoostTrack</i> , <i>online</i> [17]			
s^3	61.33	64.84	+3.51 [+0.85 , +13.28]
$T = 2$	67.59	67.59	0 (identical)

HOTA rose by 0.61 points ([+0.27, +1.65]), its MOTA by 1.13 points ([+0.27, +3.03]), and its IDF1 by 0.82 points ([+0.43, +1.94]), with higher HOTA on 7 of 7 sequences (Fig. 4). False negatives fell by 1,291, and false positives rose by 681. PD-SORT is a single-stage tracker, and its gain follows the same mechanism as for OC-SORT; unlike OC-SORT, it was not seen during design. Hybrid-SORT [18], a two-stage tracker, produced identical output, as predicted before the run.

6.5 Latency

Table 11 and Fig. 6 compare the per-frame time of the host pipeline with and without SCL. SCL took 2.95 ms per frame on average and 4.39 ms at the 95th percentile, including the phase-correlation motion cue. The end-to-end time rose from 57.85 to 60.44 ms per frame, by 4.5%. The tracker became faster, from 2.04 to 1.43 ms per frame, because it received fewer candidates. The detector takes most of each frame, about 42 ms, so the layer adds only a small part of the total latency.

7 Discussion

When the layer acts. The direction of the effect follows two properties of the pipeline rather than




Fig. 5 Detector confidence shift on MOT17 val-half: HOTA of each published tracker with its published thresholds, alone and with SCL, after four monotone transforms of the detector scores (pow3: s^3 ; scale05: $0.5s$; temp2 and temp05: $T = 2$ and $T = 0.5$).

Table 10 Baseline comparison on the predeclared external systems (MOT17 val-half, authors’ released detections), each evaluated once with the frozen SCL. Paper: number reported by the authors (for PD-SORT, their released output). Host: published tracker alone (our reproduction). Reproduction class against the authors’ number (EXACT: $|\Delta\text{HOTA}| \leq 0.2$; CLOSE: ≤ 1.0 with a documented, untuned cause). Δ : host with SCL minus host, with 95% paired-bootstrap interval; bold intervals exclude zero.

Metric	Paper	Host	Δ [95% CI]
<i>PD-SORT</i> [20], <i>EXACT</i>			
HOTA	68.01	68.01	+0.61 [+0.27 , +1.65]
MOTA	75.19	75.19	+1.13 [+0.27 , +3.03]
IDF1	81.03	81.03	+0.82 [+0.43 , +1.94]
<i>Hybrid-SORT</i> [18], <i>CLOSE</i>			
HOTA	67.1	66.70	0 (identical)
MOTA	75.8	75.52	0 (identical)
IDF1	78.0	77.56	0 (identical)

the dataset: whether the tracker has a continuation stage, and whether the detector’s score distribution matches the tracker’s thresholds. On MOT17 with well-matched detections, the two-stage trackers BoostTrack, ByteTrack at floor 0.1, and Hybrid-SORT produced output identical to their hosts, while the single-stage OC-SORT and PD-SORT gained HOTA on every sequence. When the stream fits the tracker, the layer stays out of the way, and when the tracker lacks a stage, the layer supplies it.

Self-limiting by design. In the clean regime the layer keeps the tracker’s thresholds and scores, so a tracker with a low-score stage and well-matched detections receives its native input. This is the main difference from the prior design AOB, which lowered the HOTA of the same strong trackers

Table 11 Baseline comparison of the per-frame time on a 4-vCPU Intel Xeon (2.10 GHz, no GPU) with a live YOLOv8n detector (736 pixels) and ByteTrack on three KITTI sequences (2,279 timed frames per arm). Host alone: the same pipeline without SCL. Mean / 95th percentile in milliseconds.

Stage	Host alone	Host + SCL
Frame decode	13.88 / 16.21	13.88 / 16.51
Detector (YOLOv8n, 736 px)	41.91 / 58.41	42.17 / 59.43
SCL	–	2.95 / 4.39
Tracker (ByteTrack)	2.04 / 3.57	1.43 / 2.64
End to end	57.85 / 76.35	60.44 / 80.78




Fig. 6 Mean per-frame time by stage for the host pipeline alone and with SCL (4-vCPU Intel Xeon, no GPU). The labels give the measured end-to-end time.

by 4.15 to 5.89 points (Table 4). The layer therefore does not need to know in advance whether a tracker needs help; it reads the answer from the score stream.

A continuation stage for single-stage trackers. Single-stage trackers drop every detection below one threshold, so a track can be lost when its object scores low for a few frames. The continuation rescue considers only candidates that overlap

a track of the previous frame that no usable candidate covers, so it targets exactly these gaps. OC-SORT and the external PD-SORT gained HOTA on 7 of 7 MOT17 sequences through this stage.

Detector replacement. When a detector is replaced or its confidence scale shifts, the fixed thresholds of a tracker stop matching the scores. The layer re-anchors them from the score stream without labels. Under halved scores it restored two trackers from 0 to about 66 HOTA, and with a different detector it raised the MOTA of two two-stage trackers by 6.0 and 7.9 points. This is the situation the layer was built for: a detector that is swapped, retrained, or quantized while the tracker stays fixed.

Real-time deployment. The layer needs no training, no GPU, and no access to the detector’s internals, and it needs only four contract values from the tracker’s configuration. At run time it reads the candidate list, one small grayscale image for the motion cue, and the tracks of the previous frame. On a four-core CPU it takes 2.95 ms per frame, adds 4.5% to the end-to-end time, and makes the tracker itself faster. SCL therefore adds very little latency and is suitable for real-time use. New trackers are added through their published contract values and an adapter, without changes to the policy.

8 Conclusion

This paper presented SCL, a training-free control layer that calibrates itself from a detector’s score stream and adapts a frozen tracker through a declared host contract. The layer splits the logits of recent frames with nested Otsu thresholds, estimates whether the stream is clean or noisy, keeps the tracker’s operating point on clean streams, and selects candidates, remaps scores, and sets thresholds on noisy streams. A continuation rescue gives single-stage trackers the stage they lack.

SCL was developed from always-on banding (AOB), an earlier design that intervened in every frame, and each step was compared with that baseline. One frozen configuration was then evaluated on seven published trackers, four detection sources, and two datasets with development and predeclared external roles. It raised the HOTA of

OC-SORT and of the predeclared external PD-SORT by 0.61 points, restored trackers whose thresholds no longer matched a shifted confidence scale, raised the MOTA of two-stage trackers with a noisy detector by 6.0 and 7.9 points, and left well-matched two-stage pipelines at their own accuracy. It costs 2.95 ms per frame on a CPU. These results support placing a self-calibrating layer between the detector and the tracker in pipelines where the detector changes while the tracker stays fixed.

Declarations

Funding. This research received no external funding.

Conflict of interest. The authors declare no competing interests.

Ethics approval and consent to participate. This computational study uses public benchmark data and does not involve human participants or animals.

Consent for publication. Not applicable.

Data availability. All datasets used (MOT17 and KITTI tracking) are publicly available from their providers. Tracker outputs, bootstrap records, and the frozen configuration are available from the corresponding author on reasonable request.

Code availability. The implementation of SCL and the evaluation scripts are available from the corresponding author on reasonable request.

Author contributions. Ahmed Gouda Ismail: Conceptualization, Methodology, Software, Investigation, Formal analysis, Visualization, Writing – original draft. Mohamed S. Mohamed and Tarek Ahmed Mahmoud: Supervision, Methodology, Writing – review & editing. All authors read and approved the final manuscript.

Acknowledgements. The authors thank the authors of the evaluated trackers for releasing their code, weights, and detections, and the maintainers of MOT17, KITTI, and TrackEval.

AI use in manuscript preparation. Generative AI tools were used to assist with drafting and language editing. The authors reviewed and verified all content, numerical values, and figures and take full responsibility for the manuscript.

References

- [1] Zhu, P., Wen, L., Du, D., Bian, X., Fan, H., Hu, Q., Ling, H.: Detection and Tracking Meet Drones Challenge. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **44**(11), 7380–7399 (2022) <https://doi.org/10.1109/tpami.2021.3119563>
- [2] Wu, X., Li, W., Hong, D., Tao, R., Du, Q.: Deep Learning for Unmanned Aerial Vehicle-Based Object Detection and Tracking: A survey. *IEEE Geoscience and Remote Sensing Magazine* **10**(1), 91–124 (2022) <https://doi.org/10.1109/mgrs.2021.3115137>
- [3] Lu, Y.-R., Sun, D.: Development of Real-Time Unmanned Aerial Vehicle Urban Object Detection System with Federated Learning. *Journal of Aerospace Information Systems* **21**(7), 547–553 (2024) <https://doi.org/10.2514/1.i011378>
- [4] Bewley, A., Ge, Z., Ott, L., Ramos, F., Upcroft, B.: Simple online and realtime tracking. In: 2016 IEEE International Conference on Image Processing (ICIP), pp. 3464–3468 (2016). <https://doi.org/10.1109/icip.2016.7533003>
- [5] Zhang, Y., Sun, P., Jiang, Y., Yu, D., Weng, F., Yuan, Z., Luo, P., Liu, W., Wang, X.: ByteTrack: Multi-object Tracking by Associating Every Detection Box. In: *Computer Vision – ECCV 2022. Lecture Notes in Computer Science*, pp. 1–21 (2022). https://doi.org/10.1007/978-3-031-20047-2_1
- [6] Cao, J., Pang, J., Weng, X., Khirodkar, R., Kitani, K.: Observation-Centric SORT: Rethinking SORT for Robust Multi-Object Tracking. In: 2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 9686–9696 (2023). <https://doi.org/10.1109/cvpr52729.2023.00934>
- [7] Küppers, F., Kronenberger, J., Shantia, A., Haselhoff, A.: Multivariate Confidence Calibration for Object Detection. In: 2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW), pp. 1322–1330 (2020). <https://doi.org/10.1109/cvprw50498.2020.00171>
- [8] Luiten, J., Ošep, A., Dendorfer, P., Torr, P., Geiger, A., Leal-Taixé, L., Leibe, B.: HOTA: A Higher Order Metric for Evaluating Multi-object Tracking. *International Journal of Computer Vision* **129**(2), 548–578 (2021) <https://doi.org/10.1007/s11263-020-01375-2>
- [9] Ma, L.V., Hussain, M.I., Park, J., Kim, J., Jeon, M.: Adaptive Confidence Threshold for ByteTrack in Multi-Object Tracking. In: 2023 12th International Conference on Control, Automation and Information Sciences (ICCAIS), pp. 370–374 (2023). <https://doi.org/10.1109/iccais59597.2023.10382403>
- [10] Shim, K., Ko, K., Hwang, J., Kim, C.: AdapTrack: Adaptive Thresholding-Based Matching for Multi-Object Tracking. In: 2024 IEEE International Conference on Image Processing (ICIP), pp. 2222–2228 (2024). <https://doi.org/10.1109/icip51287.2024.10648139>
- [11] Gwon, D.U., Lee, W.I.: Adaptive Threshold Control Scheme Based on Detection Confidence Distribution for Data Association in Multi-Object Tracking. *Electronics* **15**(17), 4047 (2026) <https://doi.org/10.3390/electronics15174047>
- [12] Han, Y., Huang, G., Song, S., Yang, L., Wang, H., Wang, Y.: Dynamic Neural Networks: A Survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **44**(11), 7436–7456 (2022) <https://doi.org/10.1109/tpami.2021.3117837>
- [13] Jiang, J., Ananthanarayanan, G., Bodik, P., Sen, S., Stoica, I.: Chameleon: Scalable Adaptation of Video Analytics. In: *Proceedings of the 2018 Conference of the ACM Special Interest Group on Data Communication*, pp. 253–266 (2018). <https://doi.org/10.1145/3230543.3230574>
- [14] Du, Y., Zhao, Z., Song, Y., Zhao, Y., Su, F., Gong, T., Meng, H.: StrongSORT: Make DeepSORT Great Again. *IEEE Transactions on Multimedia* **25**, 8725–8737 (2023) <https://doi.org/10.1109/tmm.2023.3240881>

- [15] Aharon, N., Orfaig, R., Bobrovsky, B.-Z.: BoT-SORT: Robust Associations Multi-Pedestrian Tracking. arXiv preprint arXiv:2206.14651 (2022)
- [16] Liu, Z., Wang, X., Wang, C., Liu, W., Bai, X.: SparseTrack: Multi-Object Tracking by Performing Scene Decomposition Based on Pseudo-Depth. *IEEE Transactions on Circuits and Systems for Video Technology* **35**(5), 4870–4882 (2025) <https://doi.org/10.1109/tcsvt.2024.3524670>
- [17] Stanojevic, V.D., Todorovic, B.T.: Boost-Track: boosting the similarity measure and detection confidence for improved multiple object tracking. *Machine Vision and Applications* **35**(3) (2024) <https://doi.org/10.1007/s00138-024-01531-5>
- [18] Yang, M., Han, G., Yan, B., Zhang, W., Qi, J., Lu, H., Wang, D.: Hybrid-SORT: Weak Cues Matter for Online Multi-Object Tracking. *Proceedings of the AAAI Conference on Artificial Intelligence* **38**(7), 6504–6512 (2024) <https://doi.org/10.1609/aaai.v38i7.28471>
- [19] Maggolino, G., Ahmad, A., Cao, J., Kitani, K.: Deep OC-Sort: Multi-Pedestrian Tracking by Adaptive Re-Identification. In: 2023 IEEE International Conference on Image Processing (ICIP), pp. 3025–3029 (2023). <https://doi.org/10.1109/icip49359.2023.10222576>
- [20] Wang, Y., Zhang, D., Li, R., Zheng, Z., Li, M.: PD-SORT: Occlusion-Robust Multi-Object Tracking Using Pseudo-Depth Cues. *IEEE Transactions on Consumer Electronics* **71**(1), 165–177 (2025) <https://doi.org/10.1109/tce.2025.3541839>
- [21] Miah, M., Bilodeau, G.-A., Saunier, N.: Learning data association for multi-object tracking using only coordinates. *Pattern Recognition* **160**, 111169 (2025) <https://doi.org/10.1016/j.patcog.2024.111169>
- [22] Cao, X., Zheng, Y., Yao, Y., Qin, H., Cao, X., Guo, S.: TOPIC: A Parallel Association Paradigm for Multi-Object Tracking Under Complex Motions and Diverse Scenes. *IEEE Transactions on Image Processing* **34**, 743–758 (2025) <https://doi.org/10.1109/tip.2025.3526066>
- [23] Shim, K., Ko, K., Yang, Y., Kim, C.: Focusing on Tracks for Online Multi-Object Tracking. In: 2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 11687–11696 (2025). <https://doi.org/10.1109/cvpr52734.2025.01091>
- [24] Zeng, F., Dong, B., Zhang, Y., Wang, T., Zhang, X., Wei, Y.: MOTR: End-to-End Multiple-Object Tracking with Transformer. In: *Computer Vision – ECCV 2022. Lecture Notes in Computer Science*, pp. 659–675 (2022). https://doi.org/10.1007/978-3-031-19812-0_38
- [25] Gao, R., Qi, J., Wang, L.: Multiple Object Tracking as ID Prediction. In: 2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 27883–27893 (2025). <https://doi.org/10.1109/cvpr52734.2025.02596>
- [26] Ge, Z., Liu, S., Wang, F., Li, Z., Sun, J.: YOLOX: Exceeding YOLO Series in 2021. arXiv preprint arXiv:2107.08430 (2021)
- [27] Jocher, G., Chaurasia, A., Qiu, J.: Ultralytics YOLOv8. <https://github.com/ultralytics/ultralytics>. Version 8.0.0 (2023)
- [28] Wang, C.-Y., Yeh, I.-H., Liao, H.-Y.M.: YOLOv9: Learning What You Want to Learn Using Programmable Gradient Information. In: *Computer Vision – ECCV 2024. Lecture Notes in Computer Science*, pp. 1–21 (2024). https://doi.org/10.1007/978-3-031-72751-1_1
- [29] Carion, N., Massa, F., Synnaeve, G., Usunier, N., Kirillov, A., Zagoruyko, S.: End-to-End Object Detection with Transformers. In: *Computer Vision – ECCV 2020. Lecture Notes in Computer Science*, pp. 213–229 (2020). https://doi.org/10.1007/978-3-030-58452-8_13
- [30] Zhao, Y., Lv, W., Xu, S., Wei, J., Wang, G.,

- Dang, Q., Liu, Y., Chen, J.: DETRs Beat YOLOs on Real-time Object Detection. In: 2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 16965–16974 (2024). <https://doi.org/10.1109/cvpr52733.2024.01605>
- [31] Otsu, N.: A Threshold Selection Method from Gray-Level Histograms. *IEEE Transactions on Systems, Man, and Cybernetics* **9**(1), 62–66 (1979) <https://doi.org/10.1109/tsmc.1979.4310076>
- [32] Reddy, B.S., Chatterji, B.N.: An FFT-based technique for translation, rotation, and scale-invariant image registration. *IEEE Transactions on Image Processing* **5**(8), 1266–1271 (1996) <https://doi.org/10.1109/83.506761>
- [33] Milan, A., Leal-Taixe, L., Reid, I., Roth, S., Schindler, K.: MOT16: A Benchmark for Multi-Object Tracking. *arXiv preprint arXiv:1603.00831* (2016)
- [34] Geiger, A., Lenz, P., Urtasun, R.: Are we ready for autonomous driving? The KITTI vision benchmark suite. In: 2012 IEEE Conference on Computer Vision and Pattern Recognition, pp. 3354–3361 (2012). <https://doi.org/10.1109/cvpr.2012.6248074>
- [35] Bernardin, K., Stiefelhagen, R.: Evaluating Multiple Object Tracking Performance: The CLEAR MOT Metrics. *EURASIP Journal on Image and Video Processing* **2008**, 1–10 (2008) <https://doi.org/10.1155/2008/246309>
- [36] Ristani, E., Solera, F., Zou, R., Cucchiara, R., Tomasi, C.: Performance Measures and a Data Set for Multi-target, Multi-camera Tracking. In: *Computer Vision – ECCV 2016 Workshops. Lecture Notes in Computer Science*, pp. 17–35 (2016). https://doi.org/10.1007/978-3-319-48881-3_2
- [37] Dendorfer, P., Ošep, A., Milan, A., Schindler, K., Cremers, D., Reid, I., Roth, S., Leal-Taixé, L.: MOTChallenge: A Benchmark for Single-Camera Multiple Target Tracking. *International Journal of Computer Vision* **129**(4), 845–881 (2021) <https://doi.org/10.1007/s11263-020-01393-0>
- [38] Efron, B., Tibshirani, R.J.: *An Introduction to the Bootstrap*, (1993). <https://doi.org/10.1007/978-1-4899-4541-9>